

Final Notebook

Please fill out:

- Student name: Andrew Levinton
- Student pace: self paced
- Scheduled project review date/time:
- Instructor name: Ahbineet Kulkarni
- Blog post URL:

Statsmodels debug

- This is because statsmodels was having version issues. this is a workaround
- The code below re-publishes the existing (but private) `_centered` function as a public attribute to the module already imported in RAM.

```
In [293]: ▶ import scipy.signal.signaltools

def _centered(arr, newsize):
    # Return the center newsize portion of the array.
    newsize = np.asarray(newsize)
    currsz = np.array(arr.shape)
    startind = (currsz - newsize) // 2
    endind = startind + newsize
    myslice = [slice(startind[k], endind[k]) for k in range(len(endind))]
    return arr[tuple(myslice)]

scipy.signal.signaltools._centered = _centered
```

Import necessary libraries

```
In [294]: # raw data handling  
import pandas as pd  
import numpy as np  
import datetime as dt  
  
# data visualiztion  
import matplotlib.pyplot as plt  
import matplotlib.cm as cm  
import seaborn as sns  
  
# regression modeling  
import statsmodels.api as sm  
from statsmodels.formula.api import ols  
from statsmodels.stats.outliers_influence import variance_inflation_factor  
  
# model validation  
from sklearn.linear_model import LinearRegression  
from sklearn.model_selection import train_test_split  
from sklearn.model_selection import cross_val_score  
from sklearn.metrics import mean_absolute_error, mean_squared_error  
  
import warnings # weird sns.distplot() warnings  
warnings.filterwarnings("ignore")  
  
plt.style.use('ggplot')
```

Define Functions

In [295]:

```

# Grabbing vifs

def get_vifs(data):
    # Get a list of the column names
    cols = data.columns

    # Create an empty DataFrame to hold the VIF results
    vif_data = pd.DataFrame()

    # Loop through each column and calculate the VIF
    for i in range(len(cols)):
        vif = variance_inflation_factor(data[cols].values, i)
        vif_data = vif_data.append({'Variable': cols[i], 'VIF': vif}, ignore_index=True)

    # Print the VIF results
    return print(vif_data)

# get ols model and plot residual distribution
def get_OLS_model(name, X, y):
    model = sm.OLS(y, sm.add_constant(X))
    results = model.fit()
    model_residual = results.resid

    return print(results.summary()), plt.suptitle(f'Residual distribution for {name}')

# get qq and histogram plots
def plot_hist_qq(df, target_col):
    """
    Creates a histogram and QQ-plot for a given dataframe and target column.

    Args:
        df (pandas.DataFrame): The dataframe to plot.
        target_col (str): The name of the target column.

    Returns:
        None
    """
    # Create subplots with 1 row and 2 columns
    fig, axs = plt.subplots(1, 2, figsize=(10, 5))

    # Plot histogram on the first subplot
    axs[0].hist(df[target_col], bins=30)
    axs[0].set_xlabel(target_col)
    axs[0].set_ylabel('Frequency')

    # Plot QQ-plot on the second subplot
    stats.probplot(df[target_col], plot=axs[1])
    axs[1].set_xlabel('Theoretical quantiles')
    axs[1].set_ylabel('Sample quantiles')

    # Adjust the layout and display the plot
    plt.tight_layout()
    plt.show()

```

```
# getting qqplots from stats model
def get_model_qqplots(data, y):
    # Set up the plot grid
    fig, axes = plt.subplots(nrows=4, ncols=6, figsize=(25, 18))

    # Loop through each variable in the DataFrame
    for i, var in enumerate(data.columns):
        # Fit a linear regression model
        X = sm.add_constant(data[var])
        model = sm.OLS(y, X).fit()

        # Calculate the residuals
        resid = model.resid

        # Create a QQ plot
        sm.qqplot(resid, line='s', ax=axes[i//6, i%6])
        axes[i//6, i%6].set_title(var)

    plt.tight_layout()
    plt.show()
```

Read in dataset, check length

In [296]: `cd data`

```
[WinError 2] The system cannot find the file specified: 'data'
C:\Users\alevi\Documents\Flatiron\dsc-data-science-env-config\Course_Folder\Phase_2\Housing_Linear_Model_Project\data
```

In [297]: `df = pd.read_csv('kc_house_data.csv')`
`len(df)`

Out[297]: 30155

Dataset timeline

In [298]: `df['yr_built'].min(), df['yr_built'].max()`

Out[298]: (1900, 2022)

Checking dtypes

In [299]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30155 entries, 0 to 30154
Data columns (total 25 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    30155 non-null  int64
1   date                 30155 non-null  object
2   price               30155 non-null  float64
3   bedrooms            30155 non-null  int64
4   bathrooms           30155 non-null  float64
5   sqft_living         30155 non-null  int64
6   sqft_lot            30155 non-null  int64
7   floors              30155 non-null  float64
8   waterfront          30155 non-null  object
9   greenbelt           30155 non-null  object
10  nuisance             30155 non-null  object
11  view                 30155 non-null  object
12  condition            30155 non-null  object
13  grade               30155 non-null  object
14  heat_source         30123 non-null  object
15  sewer_system        30141 non-null  object
16  sqft_above          30155 non-null  int64
17  sqft_basement       30155 non-null  int64
18  sqft_garage         30155 non-null  int64
19  sqft_patio          30155 non-null  int64
20  yr_built            30155 non-null  int64
21  yr_renovated        30155 non-null  int64
22  address             30155 non-null  object
23  lat                 30155 non-null  float64
24  long                30155 non-null  float64
dtypes: float64(5), int64(10), object(10)
memory usage: 5.8+ MB
```

Linear Model must meet the following assumptions:

Simple Linear Regression on select features

Assumption check:

- Is it linear?
- Is it normal?
 - histogram
 - QQ-plot
- Is it homoscedastic?

The process for building this linear model:

- Prep data for linear model regression: This involves dropping null values, dropping "bad data", as well as engineering features to assist in assuming linearization
- Key scores to look at:
 - R-Squared (or the coefficient of determination) - a statistical measure in a regression model that determines the proportion of variance in the dependent variable that can be explained by the independent variable. In other words, r-squared shows how well the data fit the regression model (the goodness of fit).
- Correlation coefficients - check to see what variables seem relatable to the target variable (price)
- residual plots - check how far data compares to the mean. Data should be normally distributed to avoid skewness of the mean
- variance inflation factor - level of statistical skew
- Root mean squared error - how far predictions fall from measured true values using Euclidean distance.
- pvalues of independent variables - measures how statistically significant the independent variables are

Data Preparation

Dropping nulls

```
In [300]: df.dropna(inplace=True)
```

Recheck length

```
In [301]: len(df)
```

```
Out[301]: 30111
```

Looking at Washington state

```
In [302]: df['address'] = df['address'].str.lower()
```

```
In [303]: df = df[df['address'].str.contains('washington')]
```

```
In [304]: len(df)
```

```
Out[304]: 29208
```

Grabbing Zipcodes

```
In [305]: df['zipcode'] = df['address'].apply(lambda x: x.split(',')[2].split(' ')[-1])
```

```
In [306]: df['zipcode'] = df['zipcode'].astype(str)
```

```
In [307]: df['zipcode'].unique()
```

```
Out[307]: array(['98055', '98133', '98178', '98118', '98027', '98166', '98030',
                '98023', '98019', '98144', '98031', '98092', '98103', '98006',
                '98136', '98007', '98038', '98057', '98077', '98126', '98053',
                '98039', '98107', '98008', '98155', '98168', '98199', '98004',
                '98045', '98052', '98011', '98002', '98033', '98116', '98198',
                '98125', '98001', '98112', '98034', '98056', '98059', '98005',
                '98040', '98014', '98106', '98029', '98122', '98003', '98117',
                '98042', '98119', '98065', '98022', '98072', '98058', '98108',
                '98115', '98074', '98105', '98024', '98146', '98109', '98102',
                '98028', '98188', '98177', '98075', '98010', '98148', '98047',
                '98032', '98070', '98051', '98288', '98354', '98272', '98296',
                '98271', '98050', '63090', 'seattle', '98387', '15301', '98251',
                '98223', '98338', '98224', '98372', '98663', '99202', '99403',
                '98422', '99203', '99223', '98270'], dtype=object)
```

Categorizing waterfronts

```
In [308]: duwamish = ['98168']
elliott_bay_zips = ['98119', '98104', '98129', '98132', '98127', '98125', '98195',
puget_sound = ['98071', '98083', '98013', '98070', '98031', '98131', '98063', '981
lake_union = ['98109']
ship_canal = ['00000']
lake_washington = ['98072', '98077']
lake_sammamish = ['98074', '98075', '98029']
other = ['00000']
river_slough_waterfronts = ['00000']

df['waterfront_loc'] = df['zipcode'].apply(lambda x: 'Duwamish' if x=='9816
else 'Elliot Bay' if x in elliott_bay_zips
else 'Puget Sound' if x in puget_sound
else 'Lake Union' if x in lake_union
else 'ship canal' if x in ship_canal
else 'Lake Washington' if x in lake_washington
else 'Lake Sammamish' if x in lake_sammamish
else 'other')
```



```
In [309]: df['waterfront_loc'].value_counts()
```

```
Out[309]: other                25497
Lake Sammamish             1159
Elliot Bay                  730
Puget Sound                 721
Lake Washington            589
Duwamish                   383
Lake Union                  129
Name: waterfront_loc, dtype: int64
```

Filter by state of Washington Zipcodes (assuming seattle is its own zipcode)

```
In [310]: df = df[df['zipcode'].str.startswith('98') | df['zipcode'].str.contains('se')]
```

One Hot Encoding Waterfronts

```
In [311]: waterfront_dummies = pd.get_dummies(df['waterfront_loc'], prefix='water', drop_first=True)
```

```
In [312]: waterfront_dummies
```

```
Out[312]:
```

	water_Elliot Bay	water_Lake Sammamish	water_Lake Union	water_Lake Washington	water_Puget Sound	water_other
0	0	0	0	0	0	1
1	0	0	0	0	0	1
2	0	0	0	0	0	1
3	0	0	0	0	0	1
4	0	0	0	0	0	1
...
30150	0	0	0	0	0	1
30151	0	0	0	0	0	1
30152	0	0	0	0	0	1
30153	0	0	0	0	0	1
30154	0	0	0	0	0	1

29200 rows × 6 columns

```
In [313]: len(df)
```

```
Out[313]: 29200
```

```
In [314]: ▶ len(df) == len(waterfront_dummies)
```

```
Out[314]: True
```

```
In [315]: ▶ df = pd.concat([df, waterfront_dummies], axis=1)
```

replacing seattle with seattle zipcode

```
In [316]: ▶ df['zipcode'] = df['zipcode'].apply(lambda x: '98101' if x=='seattle' else
```

recheck zipcodes

```
In [317]: ▶ df['zipcode'].unique()
```

```
Out[317]: array(['98055', '98133', '98178', '98118', '98027', '98166', '98030',  
                '98023', '98019', '98144', '98031', '98092', '98103', '98006',  
                '98136', '98007', '98038', '98057', '98077', '98126', '98053',  
                '98039', '98107', '98008', '98155', '98168', '98199', '98004',  
                '98045', '98052', '98011', '98002', '98033', '98116', '98198',  
                '98125', '98001', '98112', '98034', '98056', '98059', '98005',  
                '98040', '98014', '98106', '98029', '98122', '98003', '98117',  
                '98042', '98119', '98065', '98022', '98072', '98058', '98108',  
                '98115', '98074', '98105', '98024', '98146', '98109', '98102',  
                '98028', '98188', '98177', '98075', '98010', '98148', '98047',  
                '98032', '98070', '98051', '98288', '98354', '98272', '98296',  
                '98271', '98050', '98101', '98387', '98251', '98223', '98338',  
                '98224', '98372', '98663', '98422', '98270'], dtype=object)
```

```
In [318]: ▶ len(df['zipcode'].unique())
```

```
Out[318]: 89
```

Observing correlation matrix for possible features that can be used with the price

```
In [319]: df.corr()['price'].abs().sort_values(ascending=False)
```

```
Out[319]: price                1.000000
sqft_living             0.616741
sqft_above              0.546108
bathrooms               0.488039
sqft_patio              0.317623
lat                     0.296212
bedrooms                0.290994
sqft_garage              0.267477
sqft_basement           0.246548
floors                  0.199285
water_Lake Sammamish    0.141426
yr_built                0.105877
sqft_lot                 0.086790
yr_renovated            0.085506
long                    0.081940
water_Lake Washington   0.070383
water_Puget Sound       0.068457
water_other              0.064781
water_Lake Union        0.035352
id                      0.030237
water_Elliot Bay        0.004859
Name: price, dtype: float64
```

Observations

- At first glance, it appears that sqft_living, sqft_above and bathrooms are the strongest correlated features to the price.
- Further investigation is needed to measure the validity of the variables. They may be correlated with the price due to skewness or other factors that can make the correlation a deceptively "good" feature.

Changing categorical variables to numerical columns - this needs to be done if we want to use them in a linear model

```
In [320]: #extracting grade as an integer  
df['grade'] = df['grade'].apply(lambda x: int(str(x.split(' ')[0])))  
  
# replacing conditions with values  
cond_dict = {'Poor':1, 'Fair':2, 'Average':3, 'Good':4, 'Very Good':5}  
df.condition.replace(to_replace=cond_dict,inplace=True)  
  
#changing date to datetime object, get day and month  
df['date'] = pd.to_datetime(df['date'])  
df['month'] = df['date'].dt.month  
  
df['day_of_year'] = df['date'].dt.dayofyear
```

Recheck dtypes

In [321]: `df.dtypes`

```
Out[321]: id                                int64
date                                datetime64[ns]
price                               float64
bedrooms                           int64
bathrooms                          float64
sqft_living                         int64
sqft_lot                           int64
floors                             float64
waterfront                         object
greenbelt                         object
nuisance                          object
view                              object
condition                         int64
grade                             int64
heat_source                       object
sewer_system                     object
sqft_above                        int64
sqft_basement                    int64
sqft_garage                       int64
sqft_patio                       int64
yr_built                         int64
yr_renovated                     int64
address                          object
lat                              float64
long                             float64
zipcode                          object
waterfront_loc                   object
water_Elliot Bay                 uint8
water_Lake Sammamish             uint8
water_Lake Union                 uint8
water_Lake Washington            uint8
water_Puget Sound                uint8
water_other                      uint8
month                            int64
day_of_year                      int64
dtype: object
```

Extracting Numerical Predictors by filtering dtypes

In [322]: `df.dtypes.unique()`

```
Out[322]: array([dtype('int64'), dtype('<M8[ns]'), dtype('float64'), dtype('O'),
                dtype('uint8')], dtype=object)
```

```
In [323]: # categorizing dtypes  
numerical_types = ['int64', 'float64']  
numerical_predictors = list(df.select_dtypes(include=numerical_types))  
numerical_predictors
```

```
Out[323]: ['id',  
           'price',  
           'bedrooms',  
           'bathrooms',  
           'sqft_living',  
           'sqft_lot',  
           'floors',  
           'condition',  
           'grade',  
           'sqft_above',  
           'sqft_basement',  
           'sqft_garage',  
           'sqft_patio',  
           'yr_built',  
           'yr_renovated',  
           'lat',  
           'long',  
           'month',  
           'day_of_year']
```

Create dataframe of numerical values

```
In [324]: # df[numerical_predictors] selects only numerical columns  
df_numerical = df[numerical_predictors]
```

```
In [325]: df_numerical.columns
```

```
Out[325]: Index(['id', 'price', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot',  
                'floors', 'condition', 'grade', 'sqft_above', 'sqft_basement',  
                'sqft_garage', 'sqft_patio', 'yr_built', 'yr_renovated', 'lat', 'l  
ong',  
                'month', 'day_of_year'],  
               dtype='object')
```

```
In [326]: len(df_numerical)
```

```
Out[326]: 29200
```

```
In [327]: len(waterfront_dummies)
```

```
Out[327]: 29200
```

Dropping price to isolate predictors

```
In [328]: df_numerical = df_numerical.drop(['id', 'price'], axis=1)
```

```
In [329]: df_numerical['floors'] = df['floors'].astype(float)
```

Calculating variance inflation factor [VIF]

VIF levels:

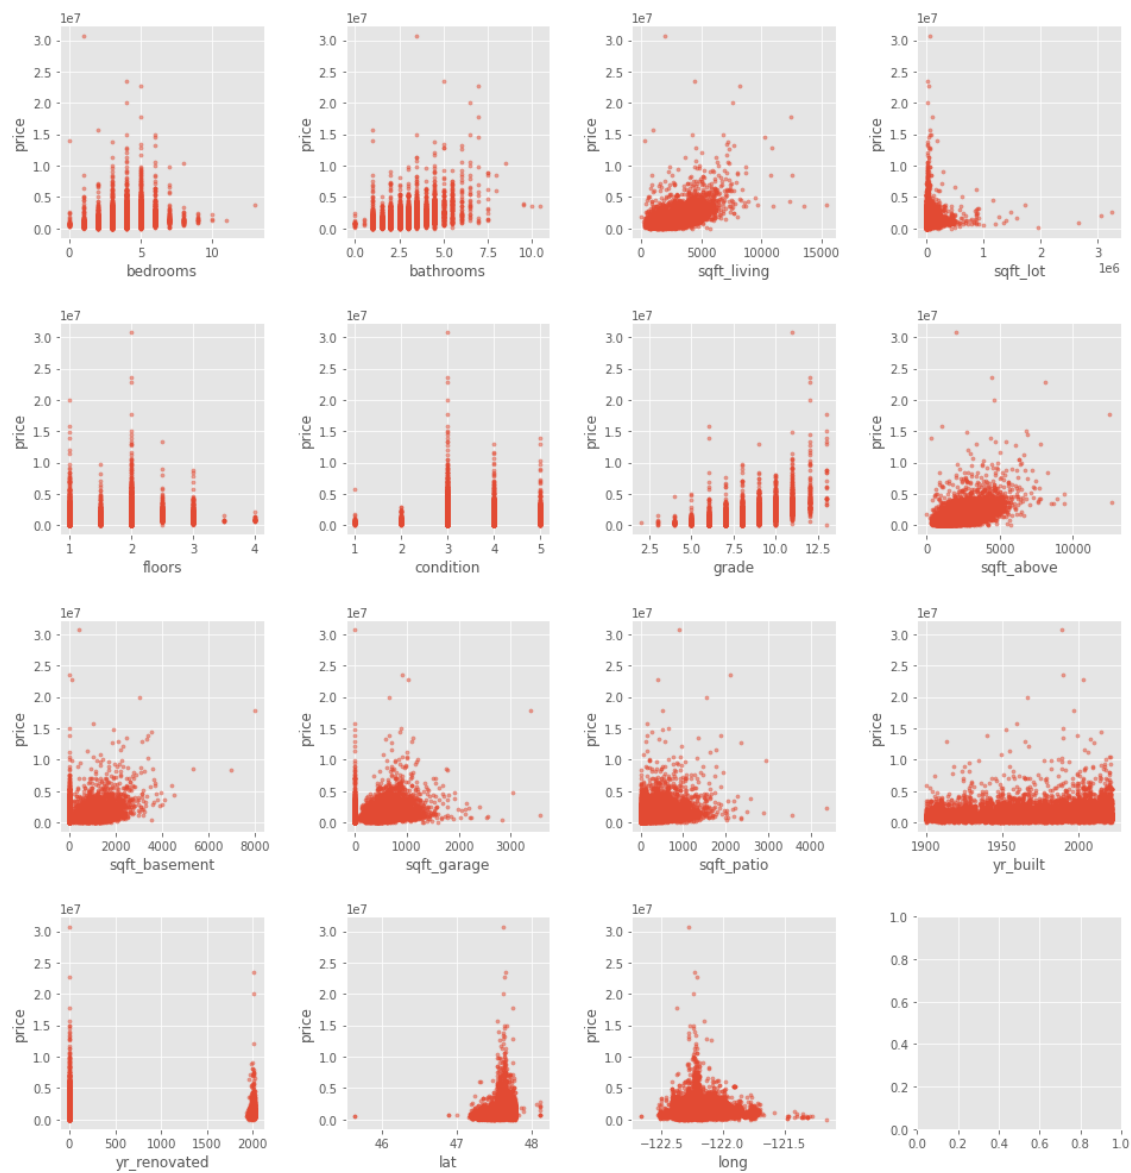
- Good: VIF ≤ 5
- Moderate/Questionable: VIF ≥ 5 and VIF ≤ 10
- Throw out: VIF ≥ 10

```
In [330]: get_vifs(df_numerical)
```

	Variable	VIF
0	bedrooms	24.749584
1	bathrooms	26.262056
2	sqft_living	119.807162
3	sqft_lot	1.140572
4	floors	17.165769
5	condition	31.148954
6	grade	131.929086
7	sqft_above	92.872396
8	sqft_basement	7.075224
9	sqft_garage	4.672471
10	sqft_patio	2.240387
11	yr_built	9220.748738
12	yr_renovated	1.210932
13	lat	110243.802192
14	long	123455.262112
15	month	697.120476
16	day_of_year	612.128232

It appears at first glance that the data only yields a small set of independent variables that are not highly collinear with each other. This will be looked at again after the removal of outliers, and the transformation of data.

```
In [331]: # Specify the dependent variable and independent variables  
y_col = 'price'  
x_cols = [col for col in df_numerical.columns if col != y_col][:15] # Use  
  
# Create scatter plot matrix  
fig, axs = plt.subplots(4, 4, figsize=(16, 16))  
for i, x_var in enumerate(x_cols):  
    row, col = divmod(i, 4)  
    axs[row, col].scatter(df_numerical[x_var], df[y_col], alpha=0.5, s=10)  
    axs[row, col].set_xlabel(x_var)  
    axs[row, col].set_ylabel(y_col)  
  
# Adjust plot layout  
fig.subplots_adjust(top=0.93, hspace=0.4, wspace=0.4)  
  
# Show the plot  
plt.show()
```

Extracting Categorical String Predictors

In [332]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 29200 entries, 0 to 30154
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     29200 non-null  int64
1   date                                 29200 non-null  datetime64[ns]
2   price                               29200 non-null  float64
3   bedrooms                           29200 non-null  int64
4   bathrooms                          29200 non-null  float64
5   sqft_living                         29200 non-null  int64
6   sqft_lot                           29200 non-null  int64
7   floors                             29200 non-null  float64
8   waterfront                         29200 non-null  object
9   greenbelt                         29200 non-null  object
10  nuisance                          29200 non-null  object
11  view                              29200 non-null  object
12  condition                         29200 non-null  int64
13  grade                             29200 non-null  int64
14  heat_source                       29200 non-null  object
15  sewer_system                     29200 non-null  object
16  sqft_above                       29200 non-null  int64
17  sqft_basement                    29200 non-null  int64
18  sqft_garage                      29200 non-null  int64
19  sqft_patio                       29200 non-null  int64
20  yr_built                         29200 non-null  int64
21  yr_renovated                     29200 non-null  int64
22  address                          29200 non-null  object
23  lat                              29200 non-null  float64
24  long                             29200 non-null  float64
25  zipcode                          29200 non-null  object
26  waterfront_loc                   29200 non-null  object
27  water_Elliot Bay                 29200 non-null  uint8
28  water_Lake Sammamish             29200 non-null  uint8
29  water_Lake Union                 29200 non-null  uint8
30  water_Lake Washington            29200 non-null  uint8
31  water_Puget Sound                29200 non-null  uint8
32  water_other                      29200 non-null  uint8
33  month                            29200 non-null  int64
34  day_of_year                      29200 non-null  int64
dtypes: datetime64[ns](1), float64(5), int64(14), object(9), uint8(6)
memory usage: 6.9+ MB
```

```
In [333]: categorical_types = ['O']  
categorical_predictors = list(df.select_dtypes(include=categorical_types))  
categorical_predictors
```

```
Out[333]: ['waterfront',  
           'greenbelt',  
           'nuisance',  
           'view',  
           'heat_source',  
           'sewer_system',  
           'address',  
           'zipcode',  
           'waterfront_loc']
```


```
In [334]: df_categorical = df[categorical_predictors]
```

In [335]: ▶ df_categorical

Out[335]:

	waterfront	greenbelt	nuisance	view	heat_source	sewer_system	address
0	NO	NO	NO	NONE	Gas	PUBLIC	2102 southeast 21st court, renton, washington ...
1	NO	NO	YES	AVERAGE	Oil	PUBLIC	11231 greenwood avenue north, seattle, washing...
2	NO	NO	NO	AVERAGE	Gas	PUBLIC	8504 south 113th street, seattle, washington 9...
3	NO	NO	NO	AVERAGE	Gas	PUBLIC	4079 letitia avenue south, seattle, washington...
4	NO	NO	YES	NONE	Electricity	PUBLIC	2193 northwest talus drive, issaquah, washingt...
...
30150	NO	NO	NO	NONE	Oil	PUBLIC	4673 eastern avenue north, seattle, washington...
30151	NO	NO	NO	FAIR	Gas	PUBLIC	4131 44th avenue southwest, seattle, washingto...
30152	NO	NO	YES	NONE	Gas	PUBLIC	910 martin luther king jr way, seattle, washin...
30153	NO	NO	NO	NONE	Gas	PUBLIC	17127 114th avenue southeast, renton, washingt...
30154	NO	NO	NO	NONE	Oil	PUBLIC	18615 7th avenue south, burien, washington 981...

29200 rows × 9 columns



Model #1

```
In [336]: ▶ model_data = df_numerical
```

```
In [337]: ▶ get_OLS_model('initial',X = model_data, y = df['price'])
```

OLS Regression Results

```

=====
=====
Dep. Variable:          price    R-squared:
0.514
Model:                  OLS      Adj. R-squared:
0.514
Method:                 Least Squares    F-statistic:
1814.
Date:                  Thu, 09 Mar 2023    Prob (F-statistic):
0.00
Time:                  20:30:28    Log-Likelihood:          -4.310
9e+05
No. Observations:      29200    AIC:                  8.62
2e+05
Df Residuals:          29182    BIC:                  8.62
4e+05
Df Model:              17
Covariance Type:       nonrobust
=====
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	-6.16e+07	4e+06	-15.396	0.000	-6.94e+07	-
5.38e+07						
bedrooms	-1.136e+05	5091.194	-22.322	0.000	-1.24e+05	-
1.04e+05						
bathrooms	9.389e+04	7527.079	12.474	0.000	7.91e+04	
1.09e+05						
sqft_living	207.5950	17.071	12.161	0.000	174.135	
241.055						
sqft_lot	0.2667	0.063	4.265	0.000	0.144	
0.389						
floors	-1.476e+05	9568.312	-15.421	0.000	-1.66e+05	-
1.29e+05						
condition	5.315e+04	5778.105	9.198	0.000	4.18e+04	
6.45e+04						
grade	2.149e+05	5521.008	38.916	0.000	2.04e+05	
2.26e+05						
sqft_above	270.4146	17.425	15.519	0.000	236.262	
304.568						
sqft_basement	80.8679	12.893	6.272	0.000	55.596	
106.140						
sqft_garage	-164.9199	18.061	-9.131	0.000	-200.320	
-129.520						
sqft_patio	193.5427	16.684	11.600	0.000	160.841	
226.244						
yr_built	-2899.2445	190.203	-15.243	0.000	-3272.051	-
2526.438						
yr_renovated	68.9239	9.331	7.386	0.000	50.634	
87.214						
lat	1.344e+06	2.68e+04	50.165	0.000	1.29e+06	
1.4e+06						
long	-1.822e+04	3.04e+04	-0.599	0.549	-7.78e+04	
4.13e+04						


```


month          1.957e+04   1.28e+04   1.529   0.126   -5515.883
4.47e+04
day_of_year    -1215.8907   420.005   -2.895   0.004   -2039.120
-392.662
=====
=====
Omnibus:                46855.092   Durbin-Watson:
1.915
Prob(Omnibus):          0.000   Jarque-Bera (JB):          9155904
1.827
Skew:                   10.060   Prob(JB):
0.00
Kurtosis:               276.586   Cond. No.                  6.9
2e+07
=====
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 6.92×10^7 . This might indicate that there are strong multicollinearity or other numerical problems.

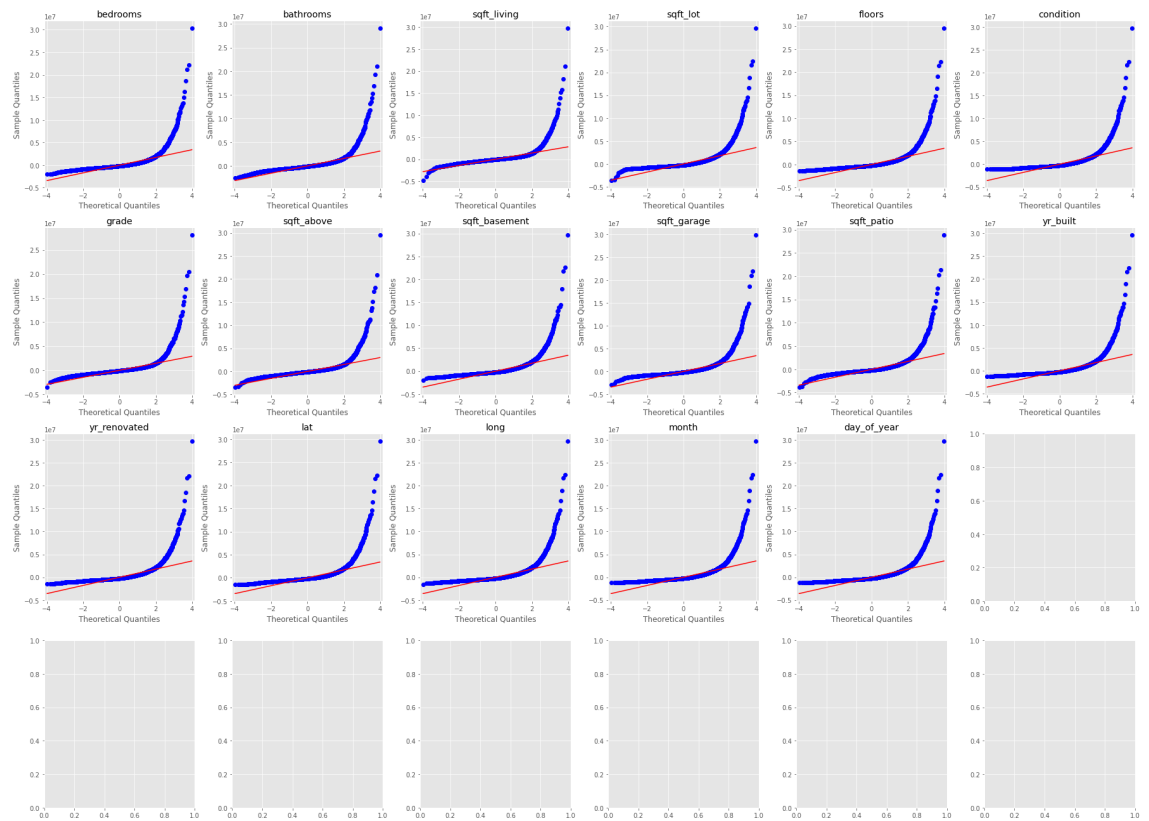


```

Out[337]: (None,
Text(0.5, 0.98, 'Residual distribution for initial model'),
<AxesSubplot:ylabel='Density'>,
None)

```

```
In [338]: get_model_qqplots(model_data, df['price'])
```



Observations

$p_value > 0.05$

- longitude **
- month
 - month was not anticipated as an effective predictor because it is not typical for the season to affect the sale price of a house

Additional Observations:

- The adjusted r-squared value is .514, indicating that this model can explain approximately 51.4% of the data.
- Skew: A kurtosis value between -2 and +2 is good to prove normalcy. The skew score is 10.065, indicating that this model is heavily skewed. This will be addressed through transformations to normalize the data.

Possible Improvements to be made to model:

- dropping of variables that are not statistically significant ($P_{val} > 0.05$)
- addition of categorical variables(one hot encoded)
- location would possibly be the most interesting variable, mapped against the waterfront or view variable
- transformation of data to satisfy normality assumption -ex: log transformation or square root transformation
- removal of outliers: Outliers in this case will be considered to be any data falling greater than 3 standard deviations outside the mean

Goals

- improve skewness - removal of outliers
- reduce homoscedacity - reduce value of VIFs
- increase rsquared to promote higher level explanation of data from model

Categorical data Exploratory Analysis and Engineering

The goal of this section will be to add in meaningful categorical data to the model, to be OneHotEncoded once prepped. For this, we first look at the categorical data.

In [339]: `df_categorical.columns`

Out[339]: Index(['waterfront', 'greenbelt', 'nuisance', 'view', 'heat_source', 'sewer_system', 'address', 'zipcode', 'waterfront_loc'], dtype='object')

Possible categorical variables of interest:

- waterfront - Whether the house is on a waterfront
 - Includes Duwamish, Elliott Bay, Puget Sound, Lake Union, Ship Canal, Lake Washington, Lake Sammamish, other lake, and river/slough waterfronts
- greenbelt - Whether the house is adjacent to a green belt
- nuisance - Whether the house has traffic noise or other recorded nuisances
- view - Quality of view from house
 - Includes views of Mt. Rainier, Olympics, Cascades, Territorial, Seattle Skyline, Puget Sound, Lake Washington, Lake Sammamish, small lake / river / creek, and other
- heat_source - Heat source for the house
- sewer_system - Sewer system for the house
- address - The street address

The grade and condition are already onehotencoded in the model and could be changed to a numerical variable, so this part of the analysis will focus on the string categorical variables.

The address appears to be the most interesting variable in the batch because it can be mapped against the waterfronts or the quality of view from the houses. For this, we will extrapolate features of the address to reduce and categorize the location.

```
In [340]: df['waterfront'].unique()
```

```
Out[340]: array(['NO', 'YES'], dtype=object)
```

```
In [341]: # convert waterfront into numeric boolean
waterfront_bool_dict = {'YES':1, 'NO':0, np.nan:0}
df_categorical.waterfront.replace(to_replace=waterfront_bool_dict, inplace=True)
```

```
In [342]: plt.scatter(x=df['waterfront'], y=df['price'])
```

```
Out[342]: <matplotlib.collections.PathCollection at 0x19ff0695130>
```



```
In [343]: df['nuisance'].unique()
```

```
Out[343]: array(['NO', 'YES'], dtype=object)
```

```
In [344]: # convert nuisance into numeric boolean
nuisance_bool_dict = {'YES':1, 'NO':0, np.nan:0}
df_categorical.nuisance.replace(to_replace=nuisance_bool_dict, inplace=True)
```

In [345]: `plt.scatter(x=df['nuisance'], y=df['price'])`

Out[345]: `<matplotlib.collections.PathCollection at 0x19f897a2490>`



In [346]: `# convert nuisance into numeric boolean
greenbelt_bool_dict = {'YES':1,'NO':0,np.nan:0}
df_categorical.greenbelt.replace(to_replace=greenbelt_bool_dict,inplace=True)`

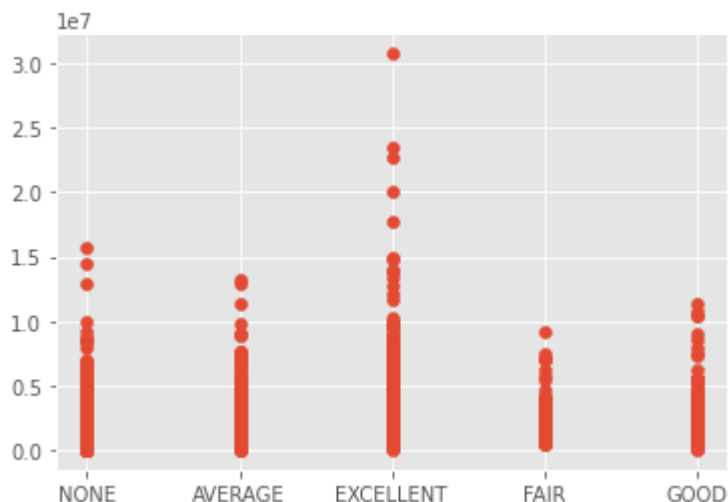
In [347]: `df['view'].unique()`

Out[347]: `array(['NONE', 'AVERAGE', 'EXCELLENT', 'FAIR', 'GOOD'], dtype=object)`

In [348]: `# convert view from string into categorical ordinal
view_dict = {'NONE':0,'FAIR':1,'AVERAGE':2,'GOOD':3,'EXCELLENT':4}
df_categorical.view.replace(to_replace=view_dict,inplace=True)`

In [349]: `plt.scatter(x=df['view'], y=df['price'])`

Out[349]: `<matplotlib.collections.PathCollection at 0x19f89a918b0>`



In [350]: `df['heat_source'].unique()`

Out[350]: array(['Gas', 'Oil', 'Electricity', 'Gas/Solar', 'Electricity/Solar',
'Other', 'Oil/Solar'], dtype=object)

In [351]: `heat_source_dummies = pd.get_dummies(df['heat_source'], prefix='heat_source')
heat_source_dummies`

Out[351]:

	heat_source_Electricity/Solar	heat_source_Gas	heat_source_Gas/Solar	heat_source_O
0	0	1	0	
1	0	0	0	
2	0	1	0	
3	0	1	0	
4	0	0	0	
...
30150	0	0	0	
30151	0	1	0	
30152	0	1	0	
30153	0	1	0	
30154	0	0	0	

29200 rows × 6 columns



In [352]: `df['sewer_system'].unique()`

Out[352]: array(['PUBLIC', 'PRIVATE', 'PRIVATE RESTRICTED', 'PUBLIC RESTRICTED'],
dtype=object)

```
In [353]: ► sewer_dummies = pd.get_dummies(df['sewer_system'], prefix='sewer', drop_first=True)
sewer_dummies
```

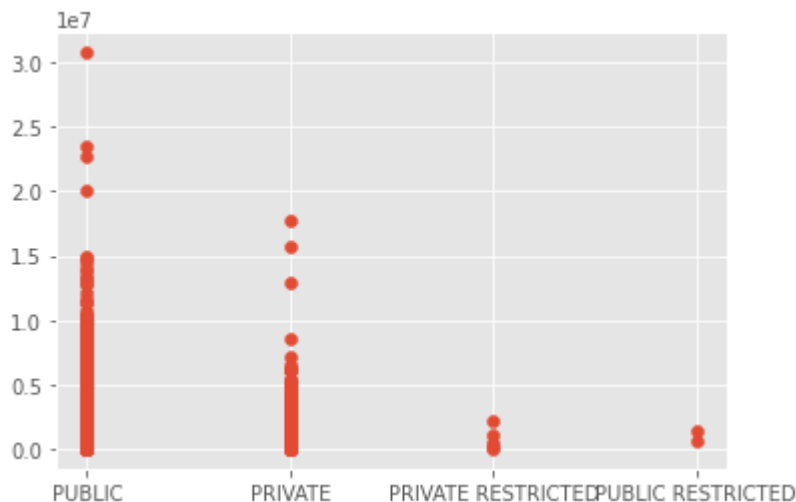
Out[353]:

	sewer_PRIVATE RESTRICTED	sewer_PUBLIC	sewer_PUBLIC RESTRICTED
0	0	1	0
1	0	1	0
2	0	1	0
3	0	1	0
4	0	1	0
...
30150	0	1	0
30151	0	1	0
30152	0	1	0
30153	0	1	0
30154	0	1	0

29200 rows × 3 columns

```
In [354]: ► plt.scatter(x=df['sewer_system'], y=df['price'])
```

Out[354]: <matplotlib.collections.PathCollection at 0x19f89be6ac0>



Developing categorical dataframe

```
In [355]: ► df_cat_pick = df_categorical[['waterfront', 'nuisance', 'view', 'greenbelt']]
```

Model #2

```
In [356]: model_2_data = pd.concat([df_numerical,sewer_dummies,heat_source_dummies, c
```

```
In [357]: len(model_2_data) == len(waterfront_dummies)
```

```
Out[357]: True
```

```
In [358]: model_2_data.columns
```

```
Out[358]: Index(['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors',  
                'condition', 'grade', 'sqft_above', 'sqft_basement', 'sqft_garag  
e',  
                'sqft_patio', 'yr_built', 'yr_renovated', 'lat', 'long', 'month',  
                'day_of_year', 'sewer_PRIVATE RESTRICTED', 'sewer_PUBLIC',  
                'sewer_PUBLIC RESTRICTED', 'heat_source_Electricity/Solar',  
                'heat_source_Gas', 'heat_source_Gas/Solar', 'heat_source_Oil',  
                'heat_source_Oil/Solar', 'heat_source_Other', 'waterfront', 'nuisa  
nce',  
                'view', 'greenbelt'],  
               dtype='object')
```



```
In [359]: ▶ get_OLS_model('second',model_2_data, df['price'])
```

OLS Regression Results

```

=====
=====
Dep. Variable:          price    R-squared:
0.555
Model:                  OLS      Adj. R-squared:
0.554
Method:                 Least Squares    F-statistic:
1211.
Date:                   Thu, 09 Mar 2023    Prob (F-statistic):
0.00
Time:                   20:30:50    Log-Likelihood:            -4.298
1e+05
No. Observations:      29200    AIC:                        8.59
7e+05
Df Residuals:          29169    BIC:                        8.59
9e+05
Df Model:               30
Covariance Type:       nonrobust
=====
=====

```

		coef	std err	t	P> t

const		-5.628e+07	4.01e+06	-14.019	0.000
-6.42e+07	-4.84e+07				
bedrooms		-8.788e+04	4935.110	-17.808	0.000
-9.76e+04	-7.82e+04				
bathrooms		8.013e+04	7265.462	11.030	0.000
6.59e+04	9.44e+04				
sqft_living		160.8497	16.429	9.791	0.000
128.649	193.050				
sqft_lot		0.3726	0.063	5.929	0.000
0.249	0.496				
floors		-1.604e+05	9224.844	-17.391	0.000
-1.79e+05	-1.42e+05				
condition		5.8e+04	5597.956	10.361	0.000
4.7e+04	6.9e+04				
grade		1.979e+05	5348.962	37.001	0.000
1.87e+05	2.08e+05				
sqft_above		295.4203	16.781	17.604	0.000
262.529	328.312				
sqft_basement		70.3779	12.464	5.647	0.000
45.948	94.807				
sqft_garage		-103.2025	17.501	-5.897	0.000
-137.504	-68.901				
sqft_patio		129.3402	16.331	7.920	0.000
97.331	161.350				
yr_built		-2419.1189	185.808	-13.019	0.000
-2783.310	-2054.927				
yr_renovated		37.1529	9.001	4.128	0.000
19.511	54.795				
lat		1.416e+06	2.6e+04	54.431	0.000
1.37e+06	1.47e+06				
long		6.158e+04	3.04e+04	2.027	0.043
2024.286	1.21e+05				

month		2.198e+04	1.23e+04	1.793	0.073
-2046.961	4.6e+04				
day_of_year		-1309.3687	402.226	-3.255	0.001
-2097.751	-520.987				
sewer_PRIVATE RESTRICTED		-5.43e+04	2.68e+05	-0.203	0.839
-5.8e+05	4.71e+05				
sewer_PUBLIC		1.703e+05	1.16e+04	14.642	0.000
1.48e+05	1.93e+05				
sewer_PUBLIC RESTRICTED		-6.324e+04	4.23e+05	-0.149	0.881
-8.92e+05	7.66e+05				
heat_source_Electricity/Solar		-3.918e+04	7.97e+04	-0.492	0.623
-1.95e+05	1.17e+05				
heat_source_Gas		-515.2456	9507.344	-0.054	0.957
-1.92e+04	1.81e+04				
heat_source_Gas/Solar		1.191e+05	6.27e+04	1.900	0.057
-3762.578	2.42e+05				
heat_source_Oil		-3.863e+04	1.45e+04	-2.656	0.008
-6.71e+04	-1.01e+04				
heat_source_Oil/Solar		-1.541e+05	2.99e+05	-0.515	0.607
-7.4e+05	4.32e+05				
heat_source_Other		-1.646e+04	1.34e+05	-0.123	0.902
-2.8e+05	2.47e+05				
waterfront		1.063e+06	3e+04	35.457	0.000
1e+06	1.12e+06				
nuisance		1.347e+04	9518.997	1.415	0.157
-5188.619	3.21e+04				
view		8.936e+04	4854.779	18.408	0.000
7.98e+04	9.89e+04				
greenbelt		7463.3683	2.24e+04	0.334	0.738
-3.63e+04	5.13e+04				


```

=====
=====
Omnibus:                45511.144    Durbin-Watson:
1.904
Prob(Omnibus):          0.000    Jarque-Bera (JB):      8590179
7.035
Skew:                   9.450    Prob(JB):
0.00
Kurtosis:               268.042    Cond. No.                7.2
5e+07
=====
=====

```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 7.25e+07. This might indicate that there are strong multicollinearity or other numerical problems.



```
Out[359]: (None,
Text(0.5, 0.98, 'Residual distribution for second model'),
<AxesSubplot:ylabel='Density'>,
None)
```

heat_source, greenbelt and sewer_system both have incredibly high p-values. These will be dropped from the final model if it holds.

Observations of Model 2

Model is still highly skewed although did present itself with some improvements. Next steps will be to normalize the data by transforming features that are skewed within the data, as well as remove outliers

- Jarque-Bera score is sky high and must come down for the model to hold any validity.
- Durbin Watson score is in the acceptable range of 1.50-2.50
- Rsquared has 'improved' but only at the expense of the the continued flaws mentioned before.

Eliminating Outliers

To normalize the distribution, outlier removal will be the first step. An outlier will be defined as three standard deviations away from the mean of the target variable.

```
In [360]: ▶ outlier_thresh = df['price'].std()*3 # value of the prices at the third standard deviation
df_outlier_removed = df.loc[abs(df['price']) <= outlier_thresh] # slicing data to remove outliers

# assign y as the target variable
y = df_outlier_removed['price']
```

```
In [361]: model_2_data_outlier_removed = model_2_data.loc[abs(df['price']) <= outlier
```

```
In [362]: df_outlier_removed
```

Out[362]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	wa
0	7399300360	2022-05-24	675000.0	4	1.0	1180	7140	1.0	
1	8910500230	2021-12-13	920000.0	5	2.5	2770	6703	1.0	
2	1180000275	2021-09-29	311000.0	6	2.0	2880	6156	1.0	
3	1604601802	2021-12-14	775000.0	3	3.0	2160	1400	2.0	
4	8562780790	2021-08-24	592500.0	2	2.0	1120	758	2.0	
...	
30150	7834800180	2021-11-30	1555000.0	5	2.0	1910	4000	1.5	
30151	194000695	2021-06-16	1313000.0	3	2.0	2020	5800	2.0	
30152	7960100080	2022-05-27	800000.0	3	2.0	1620	3600	1.0	
30153	2781280080	2022-02-24	775000.0	3	2.5	2570	2889	2.0	
30154	9557800100	2022-04-29	500000.0	3	1.5	1200	11058	1.0	

28004 rows × 35 columns

```
In [363]: waterfront_dummies = df_outlier_removed[['water_Elliot Bay', 'water_Lake Sammamish', 'water_Lake Union', 'water_Lake Washington', 'water_Puget Sound', 'water_other', 'month_of_year']]
```

```
In [364]: df_outlier_removed.columns
```

Out[364]: Index(['id', 'date', 'price', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors', 'waterfront', 'greenbelt', 'nuisance', 'view', 'condition', 'grade', 'heat_source', 'sewer_system', 'sqft_above', 'sqft_basement', 'sqft_garage', 'sqft_patio', 'yr_built', 'yr_renovated', 'address', 'lat', 'long', 'zipcode', 'waterfront_location', 'water_Elliot Bay', 'water_Lake Sammamish', 'water_Lake Union', 'water_Lake Washington', 'water_Puget Sound', 'water_other', 'month_of_year'], dtype='object')

New look at model with removed outliers

```
In [365]: outlier_data = pd.concat([y,model_2_data_outlier_removed], axis=1)
```

```
In [366]: outlier_data = outlier_data.drop('price', axis=1)
```

```
In [367]: len(outlier_data)
```


Out[367]: 28004

```
In [368]: outlier_data
```

Out[368]:

	bedrooms	bathrooms	sqft_living	sqft_lot	floors	condition	grade	sqft_above	sqft
0	4	1.0	1180	7140	1.0	4	7	1180	
1	5	2.5	2770	6703	1.0	3	7	1570	
2	6	2.0	2880	6156	1.0	3	7	1580	
3	3	3.0	2160	1400	2.0	3	9	1090	
4	2	2.0	1120	758	2.0	3	7	1120	
...
30150	5	2.0	1910	4000	1.5	4	8	1600	
30151	3	2.0	2020	5800	2.0	3	7	2020	
30152	3	2.0	1620	3600	1.0	3	7	940	
30153	3	2.5	2570	2889	2.0	3	8	1830	
30154	3	1.5	1200	11058	1.0	3	7	1200	

28004 rows × 30 columns



Model #3

```
In [369]: ▶ get_OLS_model('outlier_removed', outlier_data,y)
```

OLS Regression Results

```

=====
=====
Dep. Variable:          price    R-squared:
0.622
Model:                  OLS      Adj. R-squared:
0.622
Method:                 Least Squares    F-statistic:
1534.
Date:                   Thu, 09 Mar 2023    Prob (F-statistic):
0.00
Time:                   20:30:59    Log-Likelihood:            -3.934
7e+05
No. Observations:      28004    AIC:                        7.87
0e+05
Df Residuals:          27973    BIC:                        7.87
3e+05
Df Model:               30
Covariance Type:       nonrobust
=====
=====

```

		coef	std err	t	P> t
[0.025	0.975]				
const		-2.891e+07	2.1e+06	-13.763	0.000
-3.3e+07	-2.48e+07				
bedrooms		-1.261e+04	2647.400	-4.764	0.000
-1.78e+04	-7423.552				
bathrooms		3.438e+04	3907.140	8.800	0.000
2.67e+04	4.2e+04				
sqft_living		137.7492	8.974	15.350	0.000
120.160	155.338				
sqft_lot		0.3540	0.036	9.708	0.000
0.283	0.426				
floors		-2.695e+04	4927.866	-5.468	0.000
-3.66e+04	-1.73e+04				
condition		5.94e+04	2922.086	20.326	0.000
5.37e+04	6.51e+04				
grade		1.469e+05	2888.537	50.859	0.000
1.41e+05	1.53e+05				
sqft_above		98.6873	9.222	10.701	0.000
80.611	116.763				
sqft_basement		9.0654	6.729	1.347	0.178
-4.123	22.254				
sqft_garage		-14.9218	9.345	-1.597	0.110
-33.238	3.394				
sqft_patio		52.5853	8.886	5.918	0.000
35.168	70.002				
yr_built		-2128.0746	98.282	-21.653	0.000
-2320.712	-1935.437				
yr_renovated		29.5335	4.828	6.117	0.000
20.070	38.997				
lat		1.305e+06	1.35e+04	96.806	0.000
1.28e+06	1.33e+06				
long		2.434e+05	1.59e+04	15.356	0.000
2.12e+05	2.74e+05				

month		1.988e+04	6406.701	3.104	0.002
7326.324	3.24e+04				
day_of_year		-1128.1162	210.292	-5.365	0.000
-1540.298	-715.934				
sewer_PRIVATE RESTRICTED		1.742e+05	1.37e+05	1.268	0.205
-9.5e+04	4.43e+05				
sewer_PUBLIC		5.498e+04	6113.519	8.993	0.000
4.3e+04	6.7e+04				
sewer_PUBLIC RESTRICTED		-2.283e+04	2.17e+05	-0.105	0.916
-4.48e+05	4.02e+05				
heat_source_Electricity/Solar		-3.437e+04	4.12e+04	-0.834	0.404
-1.15e+05	4.64e+04				
heat_source_Gas		3.284e+04	4947.829	6.638	0.000
2.31e+04	4.25e+04				
heat_source_Gas/Solar		1.564e+05	3.38e+04	4.634	0.000
9.03e+04	2.23e+05				
heat_source_Oil		-1.553e+04	7536.189	-2.060	0.039
-3.03e+04	-756.860				
heat_source_Oil/Solar		-4.439e+04	1.53e+05	-0.290	0.772
-3.45e+05	2.56e+05				
heat_source_Other		9.011e+04	7.06e+04	1.276	0.202
-4.83e+04	2.29e+05				
waterfront		1.227e+05	1.82e+04	6.751	0.000
8.71e+04	1.58e+05				
nuisance		-2.687e+04	5007.274	-5.366	0.000
-3.67e+04	-1.71e+04				
view		6.205e+04	2654.342	23.375	0.000
5.68e+04	6.72e+04				
greenbelt		9.809e+04	1.19e+04	8.255	0.000
7.48e+04	1.21e+05				


```

=====
=====
Omnibus:                3918.983   Durbin-Watson:
2.002
Prob(Omnibus):          0.000   Jarque-Bera (JB):        2035
2.924
Skew:                   0.577   Prob(JB):
0.00
Kurtosis:               7.014   Cond. No.                6.6
0e+07
=====
=====

```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 6.6e+07. This might indicate that there are strong multicollinearity or other numerical problems.



```
Out[369]: (None,
Text(0.5, 0.98, 'Residual distribution for outlier_removed model'),
<AxesSubplot:ylabel='Density'>,
None)
```

```
In [370]: outlier_data.columns
```

```
Out[370]: Index(['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors',
'condition', 'grade', 'sqft_above', 'sqft_basement', 'sqft_garag
e',
'sqft_patio', 'yr_built', 'yr_renovated', 'lat', 'long', 'month',
'day_of_year', 'sewer_PRIVATE RESTRICTED', 'sewer_PUBLIC',
'sewer_PUBLIC RESTRICTED', 'heat_source_Electricity/Solar',
'heat_source_Gas', 'heat_source_Gas/Solar', 'heat_source_Oil',
'heat_source_Oil/Solar', 'heat_source_Other', 'waterfront', 'nuisa
nce',
'view', 'greenbelt'],
dtype='object')
```

Observations of model 3

pvalue > 0.05

- sqft_basement
- sqft_garage
- sewer_PRIVATE RESTRICTED
- sewer_PUBLIC RESTRICTED
- heat_source_Electricity/Solar
- heat_source_Oil/Solar
- heat_source_Other
- Adjusted rsquared indicates that the model explains 62.2% of the data.
- Skewness has improved dramatically to an acceptable range between -2 and 2. The removal of outliers has made this possible.
- Durbin-Watson score is still in the acceptable ranges of 1.5-2.5


- Jarque-Bera score is still very high but has been brought down by a significant factor. Still not perfect but trending in the right direction.
- Multicollinearity is possibly present in the model and likely so given the initial VIFs before the first model was built. VIFS should be revisited again to see if those variables are worth keeping.

Looking at transformations for the price.

```
In [371]: import scipy.stats as stats
fig, axs = plt.subplots(1, 2, figsize=(10, 5))

# Plot histogram on the first subplot
axs[0].hist(y, bins=30)
axs[0].set_xlabel('Price')
axs[0].set_ylabel('Frequency')
axs[0].set_title('histogram')
# Plot QQ-plot on the second subplot
stats.probplot(y, plot=axs[1])
axs[1].set_xlabel('Theoretical quantiles')
axs[1].set_ylabel('Sample quantiles')

# Adjust the layout and display the plot
plt.tight_layout()
plt.show()
```



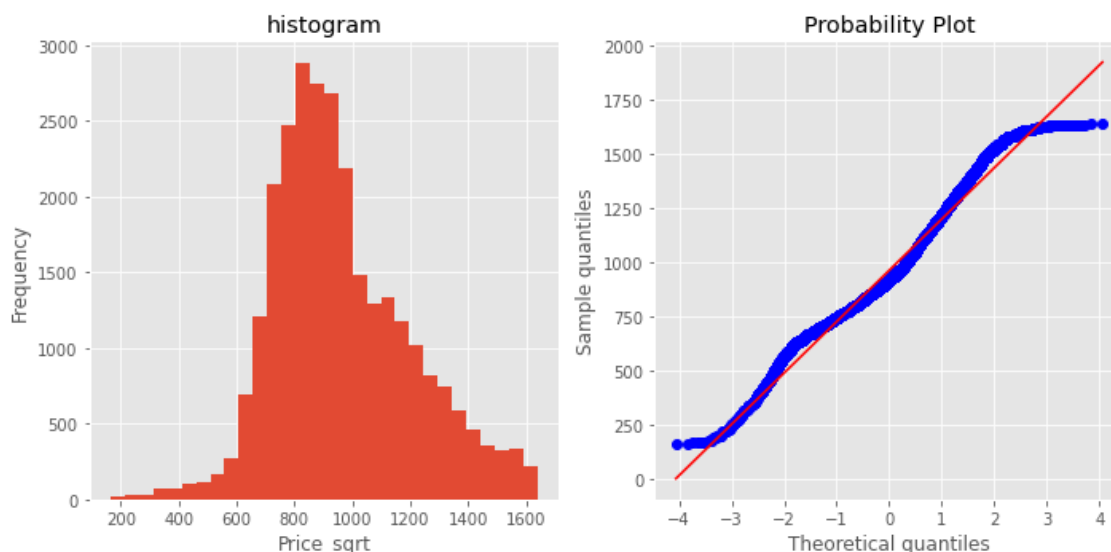
Issue above is the data shows linearization everywhere but both tails of the data. Catching the lower tail will be the goal for the next test of transformation. For this, we will try a root transformation.

```
In [372]: import matplotlib.pyplot as plt
import scipy.stats as stats

# Create subplots with 1 row and 2 columns
fig, axs = plt.subplots(1, 2, figsize=(10, 5))
y_sqrt = y**0.5
# Plot histogram on the first subplot
axs[0].hist(y_sqrt, bins=30)
axs[0].set_xlabel('Price_sqrt')
axs[0].set_ylabel('Frequency')
axs[0].set_title('histogram')

# Plot QQ-plot on the second subplot
stats.probplot(y_sqrt, plot=axs[1])
axs[1].set_xlabel('Theoretical quantiles')
axs[1].set_ylabel('Sample quantiles')

# Adjust the layout and display the plot
plt.tight_layout()
plt.show()
```



```

In [373]: fig, (ax1,ax2,ax3) = plt.subplots(1,3)


og = sns.distplot(df['price'],ax=ax1).set_title('Original data\nHeavily skewed')
ot = sns.distplot(y,ax=ax2).set_title('Removed outliers\nRemoved Skew',font
lo = sns.distplot(y_sqrt,ax=ax3).set_title('Square Root transformation\nMore

ax1.set_ylabel("")
ax2.set_ylabel("")
ax3.set_ylabel("")

plt.gcf().set_size_inches(15, 8)
plt.suptitle("Target Variable Through Iterations",fontsize=32)
fig.tight_layout()
plt.show()

```

Target Variable Through Iterations



Checking model with transformed target variable - square root transformation

```
In [374]: ▶ get_OLS_model('transformed', outlier_data, y_sqrt)
```

OLS Regression Results

```

=====
=====
Dep. Variable:          price    R-squared:
0.628
Model:                  OLS      Adj. R-squared:
0.627
Method:                 Least Squares    F-statistic:
1571.
Date:                   Thu, 09 Mar 2023    Prob (F-statistic):
0.00
Time:                   20:31:08    Log-Likelihood:            -1.794
0e+05
No. Observations:       28004    AIC:                        3.58
9e+05
Df Residuals:           27973    BIC:                        3.59
1e+05
Df Model:                30
Covariance Type:        nonrobust
=====
=====

```

		coef	std err	t	P> t

const		-1.45e+04	1005.553	-14.415	0.000
-1.65e+04	-1.25e+04				
bedrooms		-3.1025	1.267	-2.448	0.014
-5.587	-0.618				
bathrooms		19.6412	1.871	10.500	0.000
15.975	23.308				
sqft_living		0.0571	0.004	13.296	0.000
0.049	0.066				
sqft_lot		0.0002	1.75e-05	10.315	0.000
0.000	0.000				
floors		-6.1373	2.359	-2.601	0.009
-10.762	-1.513				
condition		31.1453	1.399	22.262	0.000
28.403	33.887				
grade		69.4559	1.383	50.223	0.000
66.745	72.167				
sqft_above		0.0462	0.004	10.458	0.000
0.038	0.055				
sqft_basement		0.0095	0.003	2.946	0.003
0.003	0.016				
sqft_garage		-0.0052	0.004	-1.158	0.247
-0.014	0.004				
sqft_patio		0.0268	0.004	6.293	0.000
0.018	0.035				
yr_built		-0.9368	0.047	-19.909	0.000
-1.029	-0.845				
yr_renovated		0.0138	0.002	5.970	0.000
0.009	0.018				
lat		669.3887	6.456	103.687	0.000
656.735	682.042				
long		125.8855	7.590	16.587	0.000
111.010	140.761				

month		10.4758	3.067	3.415	0.001
4.464	16.488				
day_of_year		-0.5743	0.101	-5.704	0.000
-0.772	-0.377				
sewer_PRIVATE RESTRICTED		-12.8877	65.764	-0.196	0.845
-141.788	116.013				
sewer_PUBLIC		25.4531	2.927	8.696	0.000
19.716	31.190				
sewer_PUBLIC RESTRICTED		2.5827	103.756	0.025	0.980
-200.783	205.949				
heat_source_Electricity/Solar		-37.6734	19.718	-1.911	0.056
-76.322	0.975				
heat_source_Gas		19.1408	2.369	8.080	0.000
14.498	23.784				
heat_source_Gas/Solar		63.9797	16.160	3.959	0.000
32.305	95.654				
heat_source_Oil		-1.4670	3.608	-0.407	0.684
-8.539	5.605				
heat_source_Oil/Solar		-3.7398	73.373	-0.051	0.959
-147.554	140.074				
heat_source_Other		36.1906	33.808	1.070	0.284
-30.076	102.457				
waterfront		62.8417	8.704	7.220	0.000
45.781	79.902				
nuisance		-15.3349	2.397	-6.397	0.000
-20.034	-10.636				
view		28.3706	1.271	22.325	0.000
25.880	30.861				
greenbelt		45.6410	5.689	8.022	0.000
34.490	56.792				


```

=====
=====
Omnibus:                3917.403   Durbin-Watson:
2.008
Prob(Omnibus):          0.000   Jarque-Bera (JB):        3720
1.626
Skew:                   -0.361   Prob(JB):
0.00
Kurtosis:               8.600   Cond. No.                6.6
0e+07
=====
=====

```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 6.6e+07. This might indicate that there are strong multicollinearity or other numerical problems.



```
Out[374]: (None,
Text(0.5, 0.98, 'Residual distribution for transformed model'),
<AxesSubplot:ylabel='Density'>,
None)
```

y_log vs y_sqrt

The model with the square root transformation appears to be less skewed and possesses a higher rsquared value, lending the ability of the model to explain more of the data. For these reasons we will use y_sqrt as our dependent variable for now until y_log appears to outweigh the benefit of y_sqrt.

Jarque-Beras score is significantly better as well with the y_sqrt variable so I'll go with it for now.

```
In [375]: ▶ from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error

X = outlier_data
# Fit a linear regression model
reg = LinearRegression().fit(X, y_sqrt)

# Predict the target values
y_pred = reg.predict(X)


# Calculate the mean absolute error
rmse = mean_squared_error(y_sqrt, y_pred)

print("Root mean squared error: ", rmse)
```

Root mean squared error: 21478.232538371158

Checking distribution of predictor


```
In [376]: plot_hist_qq(outlier_data, 'sqft_living')
```



Data is clearly skewed right and follows an exponential pattern similar to price. For this, we will use a logarithmic transformation.

```
In [377]: outlier_data['sqft_living_log'] = np.log(outlier_data['sqft_living'])
```

```
In [378]: plot_hist_qq(outlier_data, 'sqft_living_log')
```



```
In [379]: outlier_data = outlier_data.drop('sqft_living', axis=1)
```


In [380]:

▶ outlier_data

Out[380]:

	bedrooms	bathrooms	sqft_lot	floors	condition	grade	sqft_above	sqft_basement
0	4	1.0	7140	1.0	4	7	1180	0
1	5	2.5	6703	1.0	3	7	1570	1570
2	6	2.0	6156	1.0	3	7	1580	1580
3	3	3.0	1400	2.0	3	9	1090	1070
4	2	2.0	758	2.0	3	7	1120	550
...
30150	5	2.0	4000	1.5	4	8	1600	1130
30151	3	2.0	5800	2.0	3	7	2020	0
30152	3	2.0	3600	1.0	3	7	940	920
30153	3	2.5	2889	2.0	3	8	1830	740
30154	3	1.5	11058	1.0	3	7	1200	0

28004 rows × 30 columns



In [381]:  `get_OLS_model('transformed', outlier_data, y_sqrt)`

OLS Regression Results

```

=====
=====
Dep. Variable:          price    R-squared:
0.626
Model:                  OLS      Adj. R-squared:
0.625
Method:                 Least Squares    F-statistic:
1558.
Date:                   Thu, 09 Mar 2023    Prob (F-statistic):
0.00
Time:                   20:31:17    Log-Likelihood:            -1.794
8e+05
No. Observations:      28004    AIC:                        3.59
0e+05
Df Residuals:          27973    BIC:                        3.59
3e+05
Df Model:               30
Covariance Type:       nonrobust
=====
=====

```

```

=====
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                -1.483e+04    1012.321    -14.651    0.000
-1.68e+04    -1.28e+04
bedrooms              -1.9239         1.305     -1.474    0.141
-4.482         0.634
bathrooms             23.4883         1.859     12.634    0.000
19.844        27.132
sqft_lot              0.0002     1.75e-05     10.510    0.000
0.000         0.000
floors               -10.7543         2.335     -4.605    0.000
-15.331        -6.177
condition             32.3383         1.405     23.010    0.000
29.584        35.093
grade                70.5903         1.392     50.714    0.000
67.862        73.319
sqft_above            0.0848         0.003     26.520    0.000
0.078         0.091
sqft_basement         0.0324         0.003     11.901    0.000
0.027         0.038
sqft_garage          -0.0133         0.004     -2.997    0.003
-0.022        -0.005
sqft_patio            0.0299         0.004      7.026    0.000
0.022         0.038
yr_built              -0.8908         0.047    -18.911    0.000
-0.983        -0.798
yr_renovated          0.0144         0.002      6.188    0.000
0.010         0.019
lat                  670.4926         6.476    103.542    0.000
657.800       683.185
long                 126.1891         7.620     16.561    0.000
111.254       141.124
month                 10.3472         3.075      3.364    0.001
4.319        16.375

```

day_of_year	-0.5697	0.101	-5.644	0.000
-0.768 -0.372				
sewer_PRIVATE RESTRICTED	-1.5546	65.945	-0.024	0.981
-130.810 127.701				
sewer_PUBLIC	24.6168	2.934	8.390	0.000
18.866 30.368				
sewer_PUBLIC RESTRICTED	-7.0466	104.027	-0.068	0.946
-210.945 196.852				
heat_source_Electricity/Solar	-37.8148	19.770	-1.913	0.056
-76.566 0.936				
heat_source_Gas	18.2644	2.384	7.663	0.000
13.593 22.936				
heat_source_Gas/Solar	65.0420	16.203	4.014	0.000
33.283 96.800				
heat_source_Oil	-4.6489	3.613	-1.287	0.198
-11.731 2.433				
heat_source_Oil/Solar	-1.8443	73.570	-0.025	0.980
-146.046 142.357				
heat_source_Other	38.0187	33.898	1.122	0.262
-28.423 104.460				
waterfront	63.0894	8.727	7.229	0.000
45.984 80.195				
nuisance	-15.0994	2.404	-6.281	0.000
-19.811 -10.387				
view	29.2177	1.272	22.964	0.000
26.724 31.712				
greenbelt	47.5776	5.703	8.343	0.000
36.400 58.755				
sqft_living_log	33.6705	6.332	5.318	0.000
21.260 46.081				


```

=====
=====
Omnibus:                3862.249   Durbin-Watson:
2.007
Prob(Omnibus):          0.000   Jarque-Bera (JB):        3674
8.345
Skew:                   -0.346   Prob(JB):
0.00
Kurtosis:               8.569   Cond. No.                6.6
3e+07
=====
=====

```

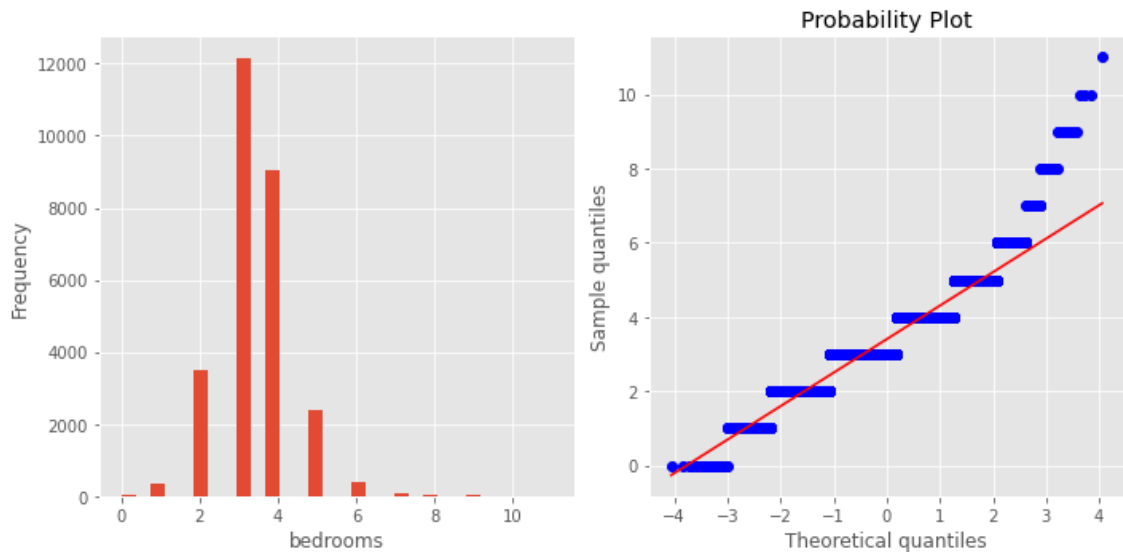
Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 6.63e+07. This might indicate that there are strong multicollinearity or other numerical problems.



```
Out[381]: (None,
Text(0.5, 0.98, 'Residual distribution for transformed model'),
<AxesSubplot:ylabel='Density'>,
None)
```

```
In [382]: plot_hist_qq(outlier_data, 'bedrooms')
```



$pval > 0.05$

- bedrooms - will be dropped from the current model

```
In [383]: outlier_data = outlier_data.drop(['bedrooms'], axis=1)
```

Rerun model

In [384]:  `get_OLS_model('transformed', outlier_data, y_sqrt)`

OLS Regression Results

```

=====
=====
Dep. Variable:          price    R-squared:
0.626
Model:                  OLS      Adj. R-squared:
0.625
Method:                 Least Squares    F-statistic:
1612.
Date:                   Thu, 09 Mar 2023    Prob (F-statistic):
0.00
Time:                   20:31:25    Log-Likelihood:            -1.794
8e+05
No. Observations:      28004    AIC:                        3.59
0e+05
Df Residuals:          27974    BIC:                        3.59
3e+05
Df Model:               29
Covariance Type:        nonrobust
=====
=====

```

		coef	std err	t	P> t

const		-1.484e+04	1012.340	-14.654	0.000
-1.68e+04	-1.29e+04				
bathrooms		22.8849	1.814	12.619	0.000
19.330	26.439				
sqft_lot		0.0002	1.75e-05	10.583	0.000
0.000	0.000				
floors		-10.5323	2.330	-4.520	0.000
-15.100	-5.965				
condition		32.2949	1.405	22.984	0.000
29.541	35.049				
grade		70.8647	1.379	51.372	0.000
68.161	73.568				
sqft_above		0.0846	0.003	26.489	0.000
0.078	0.091				
sqft_basement		0.0325	0.003	11.922	0.000
0.027	0.038				
sqft_garage		-0.0132	0.004	-2.968	0.003
-0.022	-0.004				
sqft_patio		0.0302	0.004	7.113	0.000
0.022	0.039				
yr_built		-0.8860	0.047	-18.854	0.000
-0.978	-0.794				
yr_renovated		0.0145	0.002	6.250	0.000
0.010	0.019				
lat		671.0171	6.466	103.778	0.000
658.344	683.691				
long		126.3329	7.619	16.581	0.000
111.399	141.267				
month		10.3404	3.075	3.362	0.001
4.312	16.369				
day_of_year		-0.5697	0.101	-5.643	0.000
-0.768	-0.372				

sewer_PRIVATE RESTRICTED	-1.6498	65.946	-0.025	0.980
-130.908 127.608				
sewer_PUBLIC	24.3608	2.929	8.317	0.000
18.620 30.102				
sewer_PUBLIC RESTRICTED	-8.0389	104.027	-0.077	0.938
-211.938 195.860				
heat_source_Electricity/Solar	-38.1259	19.770	-1.928	0.054
-76.876 0.624				
heat_source_Gas	18.2508	2.384	7.657	0.000
13.579 22.923				
heat_source_Gas/Solar	65.2818	16.202	4.029	0.000
33.524 97.039				
heat_source_Oil	-4.7264	3.613	-1.308	0.191
-11.808 2.355				
heat_source_Oil/Solar	-0.6240	73.567	-0.008	0.993
-144.819 143.571				
heat_source_Other	38.3315	33.898	1.131	0.258
-28.110 104.773				
waterfront	63.5322	8.722	7.284	0.000
46.437 80.628				
nuisance	-15.1146	2.404	-6.287	0.000
-19.827 -10.403				
view	29.3734	1.268	23.166	0.000
26.888 31.859				
greenbelt	47.6660	5.702	8.359	0.000
36.489 58.843				
sqft_living_log	30.9219	6.051	5.110	0.000
19.062 42.782				

```


=====
=====
Omnibus:                3860.222   Durbin-Watson:
2.007
Prob(Omnibus):          0.000   Jarque-Bera (JB):          3673
1.967
Skew:                   -0.346   Prob(JB):
0.00
Kurtosis:               8.568   Cond. No.                6.6
3e+07
=====
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 6.63e+07. This might indicate that there are strong multicollinearity or other numerical problems.



```
Out[384]: (None,  
          Text(0.5, 0.98, 'Residual distribution for transformed model'),  
          <AxesSubplot:ylabel='Density'>,  
          None)
```

Dropping sewer/heat source data

```
In [385]: new_outlier_data = outlier_data.drop(['sewer_PRIVATE RESTRICTED', 'sewer_PUE
```

In [386]:  `get_OLS_model('transformed', new_outlier_data, y_sqrt)`

OLS Regression Results

```

=====
=====
Dep. Variable:          price    R-squared:
0.626
Model:                  OLS      Adj. R-squared:
0.625
Method:                 Least Squares    F-statistic:
2032.
Date:                   Thu, 09 Mar 2023    Prob (F-statistic):
0.00
Time:                   20:31:28    Log-Likelihood:            -1.794
8e+05
No. Observations:      28004    AIC:                        3.59
0e+05
Df Residuals:          27980    BIC:                        3.59
2e+05
Df Model:               23
Covariance Type:       nonrobust
=====
=====

```

		coef	std err	t	P> t	
[0.025	0.975]					
const		-1.482e+04	1012.120	-14.646	0.000	-1.68
e+04	-1.28e+04					
bathrooms		23.1462	1.799	12.864	0.000	1
9.620	26.673					
sqft_lot		0.0002	1.75e-05	10.683	0.000	
0.000	0.000					
floors		-10.4958	2.324	-4.516	0.000	-1
5.051	-5.940					
condition		32.4584	1.393	23.299	0.000	2
9.728	35.189					
grade		70.8221	1.378	51.381	0.000	6
8.120	73.524					
sqft_above		0.0844	0.003	26.477	0.000	
0.078	0.091					
sqft_basement		0.0322	0.003	11.866	0.000	
0.027	0.038					
sqft_garage		-0.0133	0.004	-3.004	0.003	-
0.022	-0.005					
sqft_patio		0.0306	0.004	7.208	0.000	
0.022	0.039					
yr_built		-0.8757	0.046	-18.901	0.000	-
0.966	-0.785					
yr_renovated		0.0147	0.002	6.369	0.000	
0.010	0.019					
lat		671.0281	6.466	103.785	0.000	65
8.355	683.701					
long		126.5948	7.616	16.623	0.000	11
1.668	141.522					
month		10.3536	3.075	3.367	0.001	
4.327	16.380					
day_of_year		-0.5702	0.101	-5.649	0.000	-
0.768	-0.372					

sewer_PUBLIC	24.0610	2.922	8.236	0.000	1
8.334 29.788					
heat_source_Gas	19.9302	2.099	9.496	0.000	1
5.816 24.044					
heat_source_Gas/Solar	66.9906	16.163	4.145	0.000	3
5.310 98.671					
waterfront	64.3185	8.711	7.384	0.000	4
7.245 81.392					
nuisance	-15.0879	2.404	-6.276	0.000	-1
9.800 -10.376					
view	29.3481	1.267	23.159	0.000	2
6.864 31.832					
greenbelt	47.6602	5.702	8.358	0.000	3
6.483 58.837					
sqft_living_log	30.6220	6.047	5.064	0.000	1
8.770 42.474					

=====

Omnibus:	3860.717	Durbin-Watson:	
2.007			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	3668
7.399			
Skew:	-0.347	Prob(JB):	
0.00			
Kurtosis:	8.564	Cond. No.	6.6
3e+07			


=====

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 6.63e+07. This might indicate that there are strong multicollinearity or other numerical problems.



```
Out[386]: (None,
Text(0.5, 0.98, 'Residual distribution for transformed model'),
<AxesSubplot:ylabel='Density'>,
None)
```

Observations

- $pval > 0.05$

bedrooms - dropped from the current model

- all variables are statistically significant ($pvalue < 0.05$)
- Durbin-Watson Score continues to be "fine" but not improve a whole lot.
- Jarque-Bera Score continues to improve but still must come down
- skewness is now an afterthought as its at a very low -0.347 Overall no real improvement of the model happens here, we will try adding in new variables to improve as well as revisit VIFs to likely drop all that were originally at extremely high levels.

Next steps to improve the model:

1. revisit VIFs to see if any variables(now that outliers are removed and data has been transformed) should now be dropped from the model.
2. New predictors will be engineered to be added to the model. The next focus will be on the zipcodes in an attempt to narrow down the data with location-dependent price points. Possible data to be looked at are:

- waterfronts
- views
- school districts: rating, and school taxes
- tax brackets

Jarque-Beras score and skew level continue to improve but there is still some work to do.

```
In [387]: ▶ X = new_outlier_data
# Fit a linear regression model
reg = LinearRegression().fit(X, y_sqrt)

# Predict the target values
y_pred = reg.predict(X)

# Calculate the mean absolute error
rmse = mean_squared_error(y_sqrt, y_pred)

print("Root mean squared error: ", rmse)

Root mean squared error: 21598.900418299952
```

Rechecking VIFs

```
In [388]: ▶ from statsmodels.stats.outliers_influence import variance_inflation_factor

# Load your data into a pandas DataFrame
data = new_outlier_data

# Get a list of the column names
cols = data.columns

# Create an empty DataFrame to hold the VIF results
vif_data = pd.DataFrame()


# Loop through each column and calculate the VIF
for i in range(len(cols)):
    vif = variance_inflation_factor(data[cols].values, i)
    vif_data = vif_data.append({'Variable': cols[i], 'VIF': vif}, ignore_index=True)

# Print the VIF results
print(vif_data)
```

	Variable	VIF
0	bathrooms	24.288806
1	sqft_lot	1.299694
2	floors	17.333161
3	condition	31.675256
4	grade	137.636761
5	sqft_above	48.182791
6	sqft_basement	4.898542
7	sqft_garage	4.593212
8	sqft_patio	2.242563
9	yr_built	9589.111873
10	yr_renovated	1.205507
11	lat	109063.702426
12	long	123718.741939
13	month	698.983641
14	day_of_year	614.160609
15	sewer_PUBLIC	8.786042
16	heat_source_Gas	3.862553
17	heat_source_Gas/Solar	1.015119
18	waterfront	1.202680
19	nuisance	1.268623
20	view	1.425702
21	greenbelt	1.061871
22	sqft_living_log	2675.580945

Scaling data

```
In [389]: ▶ scaledX = (new_outlier_data - np.mean(new_outlier_data)) / np.std(new_outlier_data)
```


In [390]:  `get_OLS_model('scaled',scaledX, y_sqrt)`

OLS Regression Results

```

=====
=====
Dep. Variable:          price    R-squared:
0.626
Model:                  OLS      Adj. R-squared:
0.625
Method:                 Least Squares    F-statistic:
2032.
Date:                   Thu, 09 Mar 2023    Prob (F-statistic):
0.00
Time:                   20:31:40    Log-Likelihood:            -1.794
8e+05
No. Observations:      28004    AIC:                        3.59
0e+05
Df Residuals:          27980    BIC:                        3.59
2e+05
Df Model:               23
Covariance Type:       nonrobust
=====
=====

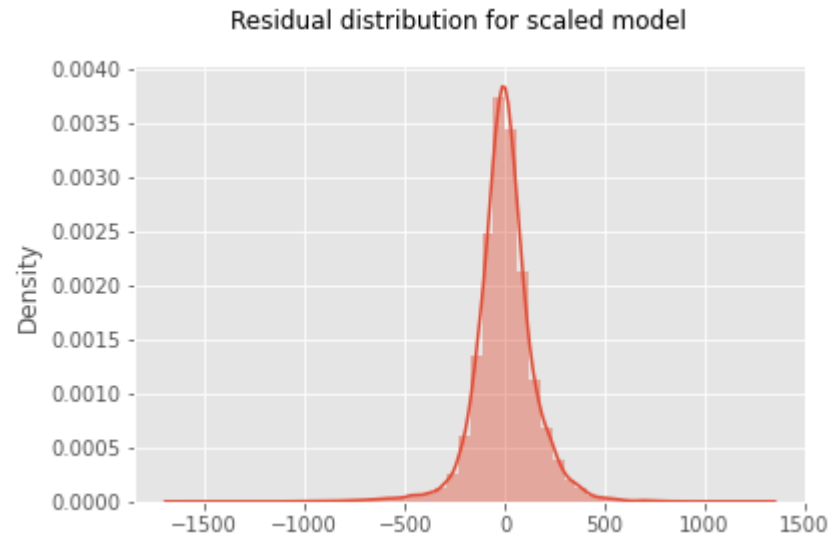
```

		coef	std err	t	P> t	
[0.025	0.975]					

const		963.8260	0.879	1097.000	0.000	96
2.104	965.548					
bathrooms		18.9843	1.476	12.864	0.000	1
6.092	21.877					
sqft_lot		10.2982	0.964	10.683	0.000	
8.409	12.188					
floors		-5.7423	1.272	-4.516	0.000	-
8.235	-3.250					
condition		23.0260	0.988	23.299	0.000	2
1.089	24.963					
grade		73.8055	1.436	51.381	0.000	7
0.990	76.621					
sqft_above		65.5334	2.475	26.477	0.000	6
0.682	70.385					
sqft_basement		17.9153	1.510	11.866	0.000	1
4.956	20.874					
sqft_garage		-3.6773	1.224	-3.004	0.003	-
6.077	-1.278					
sqft_patio		7.1036	0.985	7.208	0.000	
5.172	9.035					
yr_built		-27.5969	1.460	-18.901	0.000	-3
0.459	-24.735					
yr_renovated		6.0204	0.945	6.369	0.000	
4.168	7.873					
lat		100.0913	0.964	103.785	0.000	9
8.201	101.982					
long		18.2867	1.100	16.623	0.000	1
6.130	20.443					
month		32.0850	9.529	3.367	0.001	1
3.408	50.762					
day_of_year		-53.8286	9.528	-5.649	0.000	-7
2.505	-35.153					

```
sewer_PUBLIC      8.5826      1.042      8.236      0.000
6.540      10.625
heat_source_Gas   9.2671      0.976      9.496      0.000
7.354      11.180
heat_source_Gas/Solar  3.6635      0.884      4.145      0.000
1.931      5.396
waterfront        7.0745      0.958      7.384      0.000
5.197      8.952
nuisance          -5.6653      0.903      -6.276      0.000      -
7.435      -3.896
view             23.0870      0.997      23.159      0.000      2
1.133      25.041
greenbelt         7.5074      0.898      8.358      0.000
5.747      9.268
sqft_living_log   12.8341      2.534      5.064      0.000
7.867      17.801
=====
=====
Omnibus:          3860.717      Durbin-Watson:
2.007
Prob(Omnibus):    0.000      Jarque-Bera (JB):          3668
7.399
Skew:            -0.347      Prob(JB):
0.00
Kurtosis:        8.564      Cond. No.
33.1
=====
=====
```

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.



```
Out[390]: (None,
Text(0.5, 0.98, 'Residual distribution for scaled model'),
<AxesSubplot:ylabel='Density'>,
None)
```

Adding waterfront dummies to the model

```
In [391]: water_data = pd.concat([scaledX, waterfront_dummies], axis=1)
```

```
In [392]: water_data.columns
```

```
Out[392]: Index(['bathrooms', 'sqft_lot', 'floors', 'condition', 'grade', 'sqft_abo  
ve',  
               'sqft_basement', 'sqft_garage', 'sqft_patio', 'yr_built',  
               'yr_renovated', 'lat', 'long', 'month', 'day_of_year', 'sewer_PUBL  
IC',  
               'heat_source_Gas', 'heat_source_Gas/Solar', 'waterfront', 'nuisanc  
e',  
               'view', 'greenbelt', 'sqft_living_log', 'water_Elliot Bay',  
               'water_Lake Sammamish', 'water_Lake Washington', 'water_Puget Soun  
d',  
               'water_other'],  
              dtype='object')
```

In [393]:  `get_OLS_model('waterfront',water_data,y_sqrt)`

OLS Regression Results

```

=====
=====
Dep. Variable:          price    R-squared:
0.634
Model:                  OLS      Adj. R-squared:
0.634
Method:                 Least Squares    F-statistic:
1732.
Date:                   Thu, 09 Mar 2023    Prob (F-statistic):
0.00
Time:                   20:31:43    Log-Likelihood:            -1.791
6e+05
No. Observations:      28004    AIC:                        3.58
4e+05
Df Residuals:          27975    BIC:                        3.58
6e+05
Df Model:               28
Covariance Type:       nonrobust
=====
=====

```

		coef	std err	t	P> t	
[0.025	0.975]					

const		927.7565	6.548	141.684	0.000	91
4.922	940.591					
bathrooms		18.6287	1.459	12.766	0.000	1
5.769	21.489					
sqft_lot		11.0210	0.955	11.543	0.000	
9.150	12.892					
floors		-6.0926	1.258	-4.844	0.000	-
8.558	-3.627					
condition		23.5247	0.978	24.065	0.000	2
1.609	25.441					
grade		69.5427	1.431	48.587	0.000	6
6.737	72.348					
sqft_above		64.4651	2.448	26.330	0.000	5
9.666	69.264					
sqft_basement		17.7160	1.494	11.857	0.000	1
4.787	20.645					
sqft_garage		-3.3880	1.214	-2.791	0.005	-
5.768	-1.008					
sqft_patio		7.8678	0.975	8.070	0.000	
5.957	9.779					
yr_built		-26.1333	1.447	-18.056	0.000	-2
8.970	-23.296					
yr_renovated		6.5953	0.935	7.055	0.000	
4.763	8.428					
lat		100.4427	1.010	99.490	0.000	9
8.464	102.422					
long		11.5037	1.127	10.205	0.000	
9.294	13.713					
month		31.1044	9.420	3.302	0.001	1
2.641	49.568					
day_of_year		-52.8929	9.420	-5.615	0.000	-7
1.356	-34.430					

sewer_PUBLIC	5.4042	1.064	5.079	0.000	
3.319 7.490					
heat_source_Gas	9.3205	0.965	9.654	0.000	
7.428 11.213					
heat_source_Gas/Solar	3.5643	0.874	4.079	0.000	
1.851 5.277					
waterfront	7.0714	0.949	7.448	0.000	
5.210 8.932					
nuisance	-5.3175	0.893	-5.953	0.000	-
7.068 -3.567					
view	23.1952	0.986	23.517	0.000	2
1.262 25.128					
greenbelt	7.5679	0.888	8.520	0.000	
5.827 9.309					
sqft_living_log	14.7779	2.507	5.894	0.000	
9.864 19.692					
water_Elliot Bay	-0.4515	8.572	-0.053	0.958	-1
7.253 16.349					
water_Lake Sammamish	140.2069	8.237	17.022	0.000	12
4.062 156.352					
water_Lake Washington	-19.8540	9.433	-2.105	0.035	-3
8.343 -1.365					
water_Puget Sound	6.8364	8.548	0.800	0.424	-
9.918 23.591					
water_other	35.4111	6.609	5.358	0.000	2
2.457 48.366					

=====

=====


Omnibus:	3994.853	Durbin-Watson:	
2.002			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	4061
9.628			
Skew:	-0.349	Prob(JB):	
0.00			
Kurtosis:	8.859	Cond. No.	
41.8			

=====

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.



```
Out[393]: (None,
Text(0.5, 0.98, 'Residual distribution for waterfront model'),
<AxesSubplot:ylabel='Density'>,
None)
```

Elliot Bay and Puget Sound present high pvalues indicating a lack of statistical significance. These will be dropped from the model.

```
In [394]: water_data = water_data.drop(['water_Elliot Bay', 'water_Puget Sound'], axis=1)
```


In [395]:  `get_OLS_model('waterfront',water_data,y_sqrt)`

OLS Regression Results

```

=====
=====
Dep. Variable:          price    R-squared:
0.634
Model:                  OLS      Adj. R-squared:
0.634
Method:                 Least Squares    F-statistic:
1865.
Date:                   Thu, 09 Mar 2023    Prob (F-statistic):
0.00
Time:                   20:31:45    Log-Likelihood:            -1.791
6e+05
No. Observations:      28004    AIC:                        3.58
4e+05
Df Residuals:          27977    BIC:                        3.58
6e+05
Df Model:               26
Covariance Type:       nonrobust
=====
=====

```

		coef	std err	t	P> t	
[0.025	0.975]					
const		930.1301	3.373	275.797	0.000	92
3.520	936.740					
bathrooms		18.6098	1.459	12.755	0.000	1
5.750	21.470					
sqft_lot		11.0471	0.954	11.577	0.000	
9.177	12.917					
floors		-6.1257	1.257	-4.872	0.000	-
8.590	-3.661					
condition		23.5314	0.977	24.077	0.000	2
1.616	25.447					
grade		69.5142	1.431	48.584	0.000	6
6.710	72.319					
sqft_above		64.4517	2.448	26.332	0.000	5
9.654	69.249					
sqft_basement		17.6850	1.493	11.842	0.000	1
4.758	20.612					
sqft_garage		-3.3364	1.213	-2.751	0.006	-
5.714	-0.959					
sqft_patio		7.8816	0.975	8.087	0.000	
5.971	9.792					
yr_built		-26.0584	1.445	-18.028	0.000	-2
8.891	-23.225					
yr_renovated		6.5990	0.935	7.060	0.000	
4.767	8.431					
lat		100.3041	0.995	100.854	0.000	9
8.355	102.253					
long		11.4506	1.126	10.171	0.000	
9.244	13.657					
month		31.0770	9.420	3.299	0.001	1
2.614	49.540					
day_of_year		-52.8566	9.419	-5.611	0.000	-7
1.319	-34.394					

sewer_PUBLIC	5.3772	1.062	5.061	0.000	
3.295 7.460					
heat_source_Gas	9.3409	0.965	9.678	0.000	
7.449 11.233					
heat_source_Gas/Solar	3.5717	0.874	4.088	0.000	
1.859 5.284					
waterfront	7.1336	0.947	7.530	0.000	
5.277 8.990					
nuisance	-5.3242	0.893	-5.965	0.000	-
7.074 -3.575					
view	23.1869	0.986	23.512	0.000	2
1.254 25.120					
greenbelt	7.5681	0.888	8.521	0.000	
5.827 9.309					
sqft_living_log	14.8136	2.506	5.910	0.000	
9.901 19.726					
water_Lake Sammamish	137.8850	5.965	23.117	0.000	12
6.194 149.576					
water_Lake Washington	-22.1141	7.532	-2.936	0.003	-3
6.878 -7.350					
water_other	33.0328	3.495	9.453	0.000	2
6.183 39.882					

=====

=====

Omnibus:	3992.548	Durbin-Watson:	
2.002			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	4059
4.858			
Skew:	-0.348	Prob(JB):	
0.00			
Kurtosis:	8.857	Cond. No.	
33.1			


=====

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Residual distribution for waterfront model



```
Out[395]: (None,
          Text(0.5, 0.98, 'Residual distribution for waterfront model'),
          <AxesSubplot:ylabel='Density'>,
          None)
```

Recheck VIFs

```
In [396]: ▶ get_vifs(water_data)
```

	Variable	VIF
0	bathrooms	2.822139
1	sqft_lot	1.207074
2	floors	2.095512
3	condition	1.266178
4	grade	2.708837
5	sqft_above	7.942106
6	sqft_basement	2.955381
7	sqft_garage	1.950322
8	sqft_patio	1.259318
9	yr_built	2.769200
10	yr_renovated	1.158303
11	lat	1.311024
12	long	1.646079
13	month	117.635393
14	day_of_year	117.622422
15	sewer_PUBLIC	1.494553
16	heat_source_Gas	1.235084
17	heat_source_Gas/Solar	1.012087
18	waterfront	1.189363
19	nuisance	1.056349
20	view	1.288746
21	greenbelt	1.045751
22	sqft_living_log	8.326010
23	water_Lake Sammamish	1.133347
24	water_Lake Washington	1.160076
25	water_other	1.008041

Month and day_of_year present with high variance inflation factors indicating possible collinearity. These will be dropped.

```
In [397]: ▶ water_data = water_data.drop(['month', 'day_of_year'], axis = 1)
```

```
In [398]: get_vifs(water_data)
```

	Variable	VIF
0	bathrooms	2.820882
1	sqft_lot	1.206917
2	floors	2.095106
3	condition	1.265521
4	grade	2.708681
5	sqft_above	7.941105
6	sqft_basement	2.955260
7	sqft_garage	1.949919
8	sqft_patio	1.259183
9	yr_built	2.767781
10	yr_renovated	1.158123
11	lat	1.311003
12	long	1.645893
13	sewer_PUBLIC	1.494486
14	heat_source_Gas	1.235073
15	heat_source_Gas/Solar	1.012068
16	waterfront	1.189304
17	nuisance	1.056119
18	view	1.288704
19	greenbelt	1.045747
20	sqft_living_log	8.325701
21	water_Lake Sammamish	1.133158
22	water_Lake Washington	1.159873
23	water_other	1.007988

All VIFs are now below 10 with the majority less than 3, meaning the issue of collinearity is now for the most part solved.

Final model

```
In [399]: ▶ get_OLS_model('waterfront', water_data, y_sqrt)
```

OLS Regression Results

```

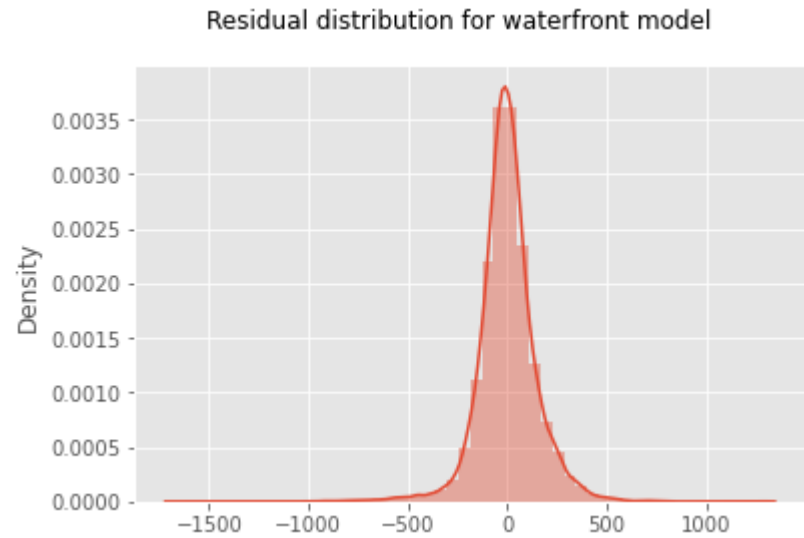
=====
=====
Dep. Variable:          price    R-squared:
0.626
Model:                  OLS      Adj. R-squared:
0.625
Method:                 Least Squares    F-statistic:
1949.
Date:                   Thu, 09 Mar 2023    Prob (F-statistic):
0.00
Time:                   20:31:59    Log-Likelihood:            -1.794
8e+05
No. Observations:       28004    AIC:                        3.59
0e+05
Df Residuals:           27979    BIC:                        3.59
2e+05
Df Model:                24
Covariance Type:        nonrobust
=====
=====

```

		coef	std err	t	P> t	
[0.025	0.975]					
const		929.1767	3.411	272.414	0.000	92
2.491	935.862					
bathrooms		17.8321	1.475	12.086	0.000	1
4.940	20.724					
sqft_lot		10.7765	0.965	11.166	0.000	
8.885	12.668					
floors		-5.7693	1.272	-4.537	0.000	-
8.262	-3.277					
condition		23.4380	0.988	23.717	0.000	2
1.501	25.375					
grade		69.7639	1.447	48.208	0.000	6
6.927	72.600					
sqft_above		63.9056	2.476	25.815	0.000	5
9.053	68.758					
sqft_basement		17.4519	1.510	11.554	0.000	1
4.491	20.412					
sqft_garage		-3.1052	1.227	-2.531	0.011	-
5.509	-0.701					
sqft_patio		7.7906	0.986	7.903	0.000	
5.858	9.723					
yr_built		-26.3626	1.462	-18.037	0.000	-2
9.227	-23.498					
yr_renovated		6.5683	0.945	6.948	0.000	
4.715	8.421					
lat		100.3684	1.006	99.778	0.000	9
8.397	102.340					
long		11.7060	1.139	10.281	0.000	
9.474	13.938					
sewer_PUBLIC		5.4272	1.075	5.050	0.000	
3.321	7.533					
heat_source_Gas		9.3428	0.976	9.570	0.000	
7.429	11.256					

```
heat_source_Gas/Solar      3.5477      0.884      4.014      0.000
1.816      5.280
waterfront      7.0093      0.958      7.315      0.000
5.131      8.887
nuisance      -5.6582      0.903      -6.268      0.000      -
7.428      -3.889
view      23.0579      0.997      23.117      0.000      2
1.103      25.013
greenbelt      7.5699      0.898      8.427      0.000
5.809      9.331
sqft_living_log      15.1773      2.535      5.987      0.000      1
0.209      20.146
water_Lake Sammamish      137.7419      6.033      22.833      0.000      12
5.918      149.566
water_Lake Washington      -22.6046      7.618      -2.967      0.003      -3
7.537      -7.672
water_other      34.1416      3.534      9.660      0.000      2
7.214      41.069
=====
=====
Omnibus:      3707.653      Durbin-Watson:
2.007
Prob(Omnibus):      0.000      Jarque-Bera (JB):      3577
0.221
Skew:      -0.301      Prob(JB):
0.00
Kurtosis:      8.504      Cond. No.
21.4
=====
=====
```

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.



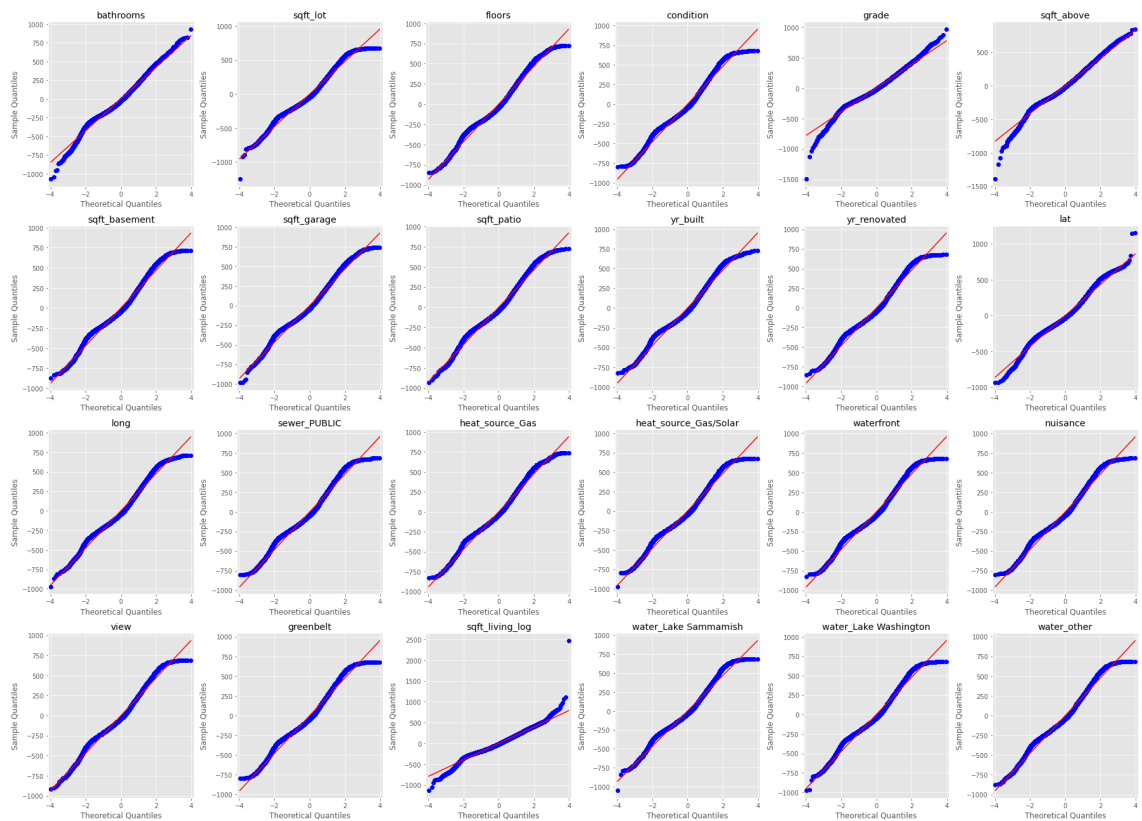
```
Out[399]: (None,
Text(0.5, 0.98, 'Residual distribution for waterfront model'),
<AxesSubplot:ylabel='Density'>,
None)
```


In [400]: `water_data.columns`

Out[400]: Index(['bathrooms', 'sqft_lot', 'floors', 'condition', 'grade', 'sqft_above',
 'sqft_basement', 'sqft_garage', 'sqft_patio', 'yr_built',
 'yr_renovated', 'lat', 'long', 'sewer_PUBLIC', 'heat_source_Gas',
 'heat_source_Gas/Solar', 'waterfront', 'nuisance', 'view', 'greenbelt',
 'sqft_living_log', 'water_Lake Sammamish', 'water_Lake Washington',
 'water_other'],
 dtype='object')

Constructing QQplots for all independent variables within the model

In [401]: `get_model_qqplots(water_data, y_sqrt)`



```
In [402]: ▶ model = sm.OLS(y_sqrt, sm.add_constant(water_data))
          results = model.fit()
          model_residual = results.resid
          model_params = results.params

          print(results.params)
```

```
const                929.176743
bathrooms            17.832111
sqft_lot             10.776486
floors               -5.769272
condition            23.438047
grade                69.763857
sqft_above           63.905587
sqft_basement        17.451898
sqft_garage          -3.105150
sqft_patio           7.790561
yr_built             -26.362631
yr_renovated          6.568336
lat                  100.368386
long                 11.706009
sewer_PUBLIC          5.427152
heat_source_Gas       9.342830
heat_source_Gas/Solar 3.547675
waterfront            7.009294
nuisance             -5.658225
view                 23.057945
greenbelt             7.569852
sqft_living_log       15.177339
water_Lake Sammamish  137.741862
water_Lake Washington -22.604616
water_other           34.141607
dtype: float64
```

We have a linear model with the dependent variable (price) square root transformed, and the following independent variables and their corresponding coefficients:

- Bathrooms: 17.832111
- Sqft_lot: 10.776486
- Floors: -5.769272
- Condition: 23.438047
- Grade: 69.763857
- Sqft_above: 63.905587
- Sqft_basement: 17.451898
- Sqft_garage: -3.105150
- Sqft_patio: 7.790561
- Yr_built: -26.362631
- Yr_renovated: 6.568336
- Lat: 100.368386
- Long: 11.706009
- Sewer_PUBLIC: 5.427152
- Heat_source_Gas: 9.342830
- Heat_source_Gas/Solar: 3.547675

- Waterfront: 7.009294
- Nuisance: -5.658225
- View: 23.057945
- Greenbelt: 7.569852
- Sqft_living_log: 15.177339
- Water_Lake Sammamish: 137.741862
- Water_Lake Washington: -22.604616
- Water_other: 34.141607

A positive coefficient indicates that as the corresponding independent variable increases, the square root of the price of the house also increases, while a negative coefficient indicates that as the corresponding independent variable increases, the square root of the price of the house decreases.

In this model, we see that the most important variable in predicting the square root of house prices is the latitude of the house, with a coefficient of 100.368386. This suggests that houses located further north tend to have higher prices. The next most important variable is water proximity, with Water_Lake Sammamish variable having a very high coefficient of 137.741862, suggesting that houses located near this lake tend to have much higher prices than other houses. On the other hand, the Water_Lake Washington variable has a negative coefficient, indicating that houses located near this lake tend to have lower prices than other houses.

Other important variables include the grade of the house, the square footage of the house above ground, and the condition of the house, all with coefficients greater than 20. The number of bathrooms, square footage of the basement, and the size of the view from the house are also important, with coefficients greater than 15.

On the other hand, variables such as the square footage of the garage and the presence of a nuisance nearby have negative coefficients, indicating that houses with larger garages or located near nuisances tend to have lower prices. The year the house was built and the longitude of the house also have negative coefficients, suggesting that older houses and houses located further west tend to have lower prices.

Overall, these results suggest that there are many factors that contribute to the price of a house, and that location, house size and quality, and the presence of nearby amenities all play important roles in determining the square root of house prices.